

Pagination done the Right Way



@MarkusWinand
@SQLPerfTips

Takeaways

- ▶ OFFSET kills performance
- ▶ Paging can be done without OFFSET
- ▶ It's faster and has fewer side effects

Note

In this presentation
index means B-tree index.

A Trivial Example

A query to fetch the 10 most recent news:

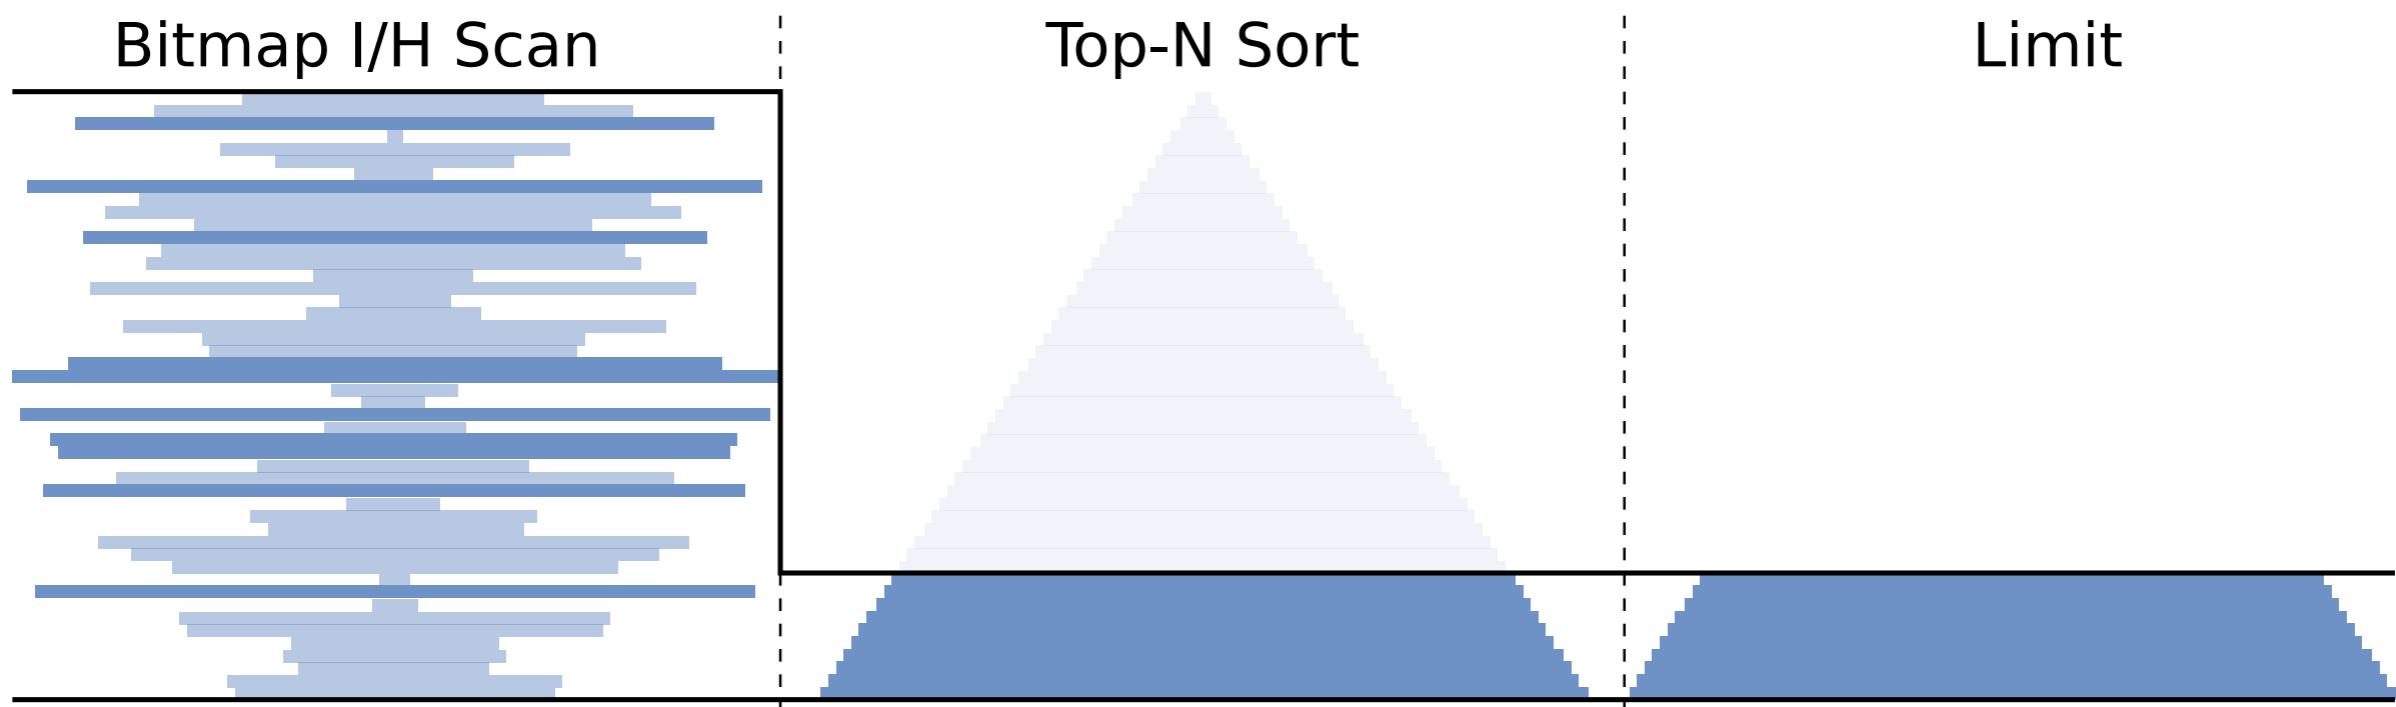
```
select *
  from news
 where topic = 1234
 order by date desc, id desc
 limit 10;

create index .. on news(topic);
```

Using `order by` to get the most recent first and
limit to fetch only the first 10.

Alternative SQL-2008 syntax (since PostgreSQL 8.4)
`fetch first 10 rows only`

Worst Case: No Index for `order by`



Limit (**actual rows=10**)

-> Sort (**actual rows=10**)

Sort Method: **top-N heapsort** Memory: 18kB

-> Bitmap Heap Scan (**rows=10000**)

Recheck Cond: (topic = 1234)

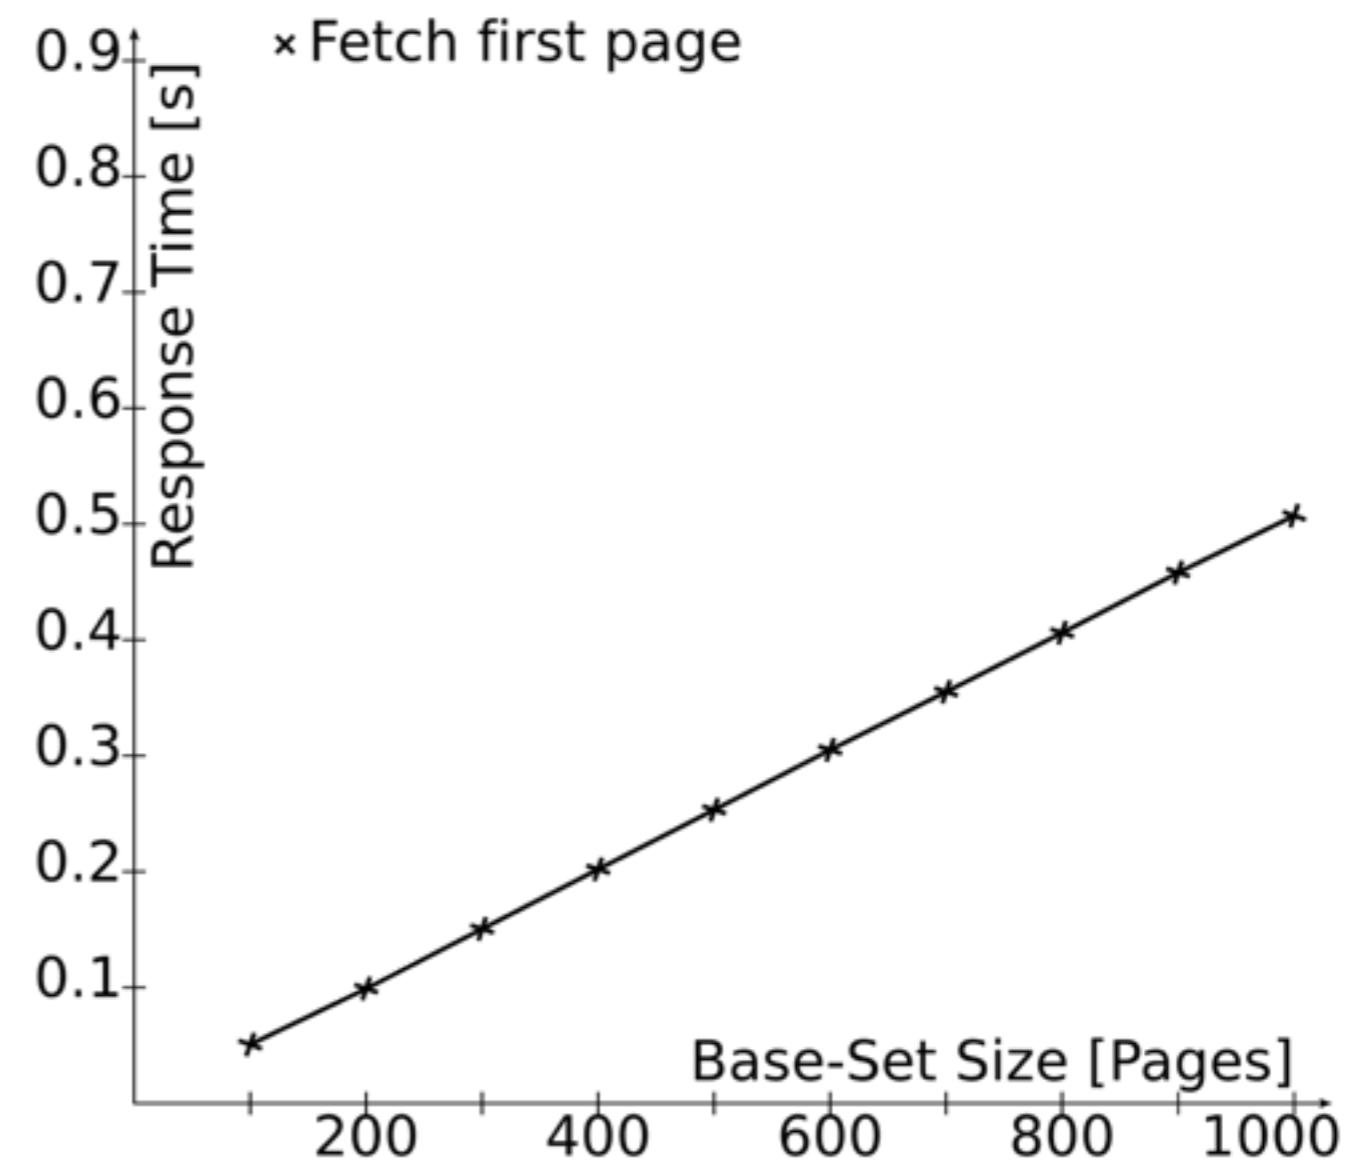
-> Bitmap Index Scan (**rows=10000**)

Index Cond: (topic = 1234)

Worst Case: No Index for order by

The limiting factor is the number of rows that match the where clause (“Base-Set Size”).

The database might use an index to satisfy the where clause, but must still fetch all matching rows to “sort” them.

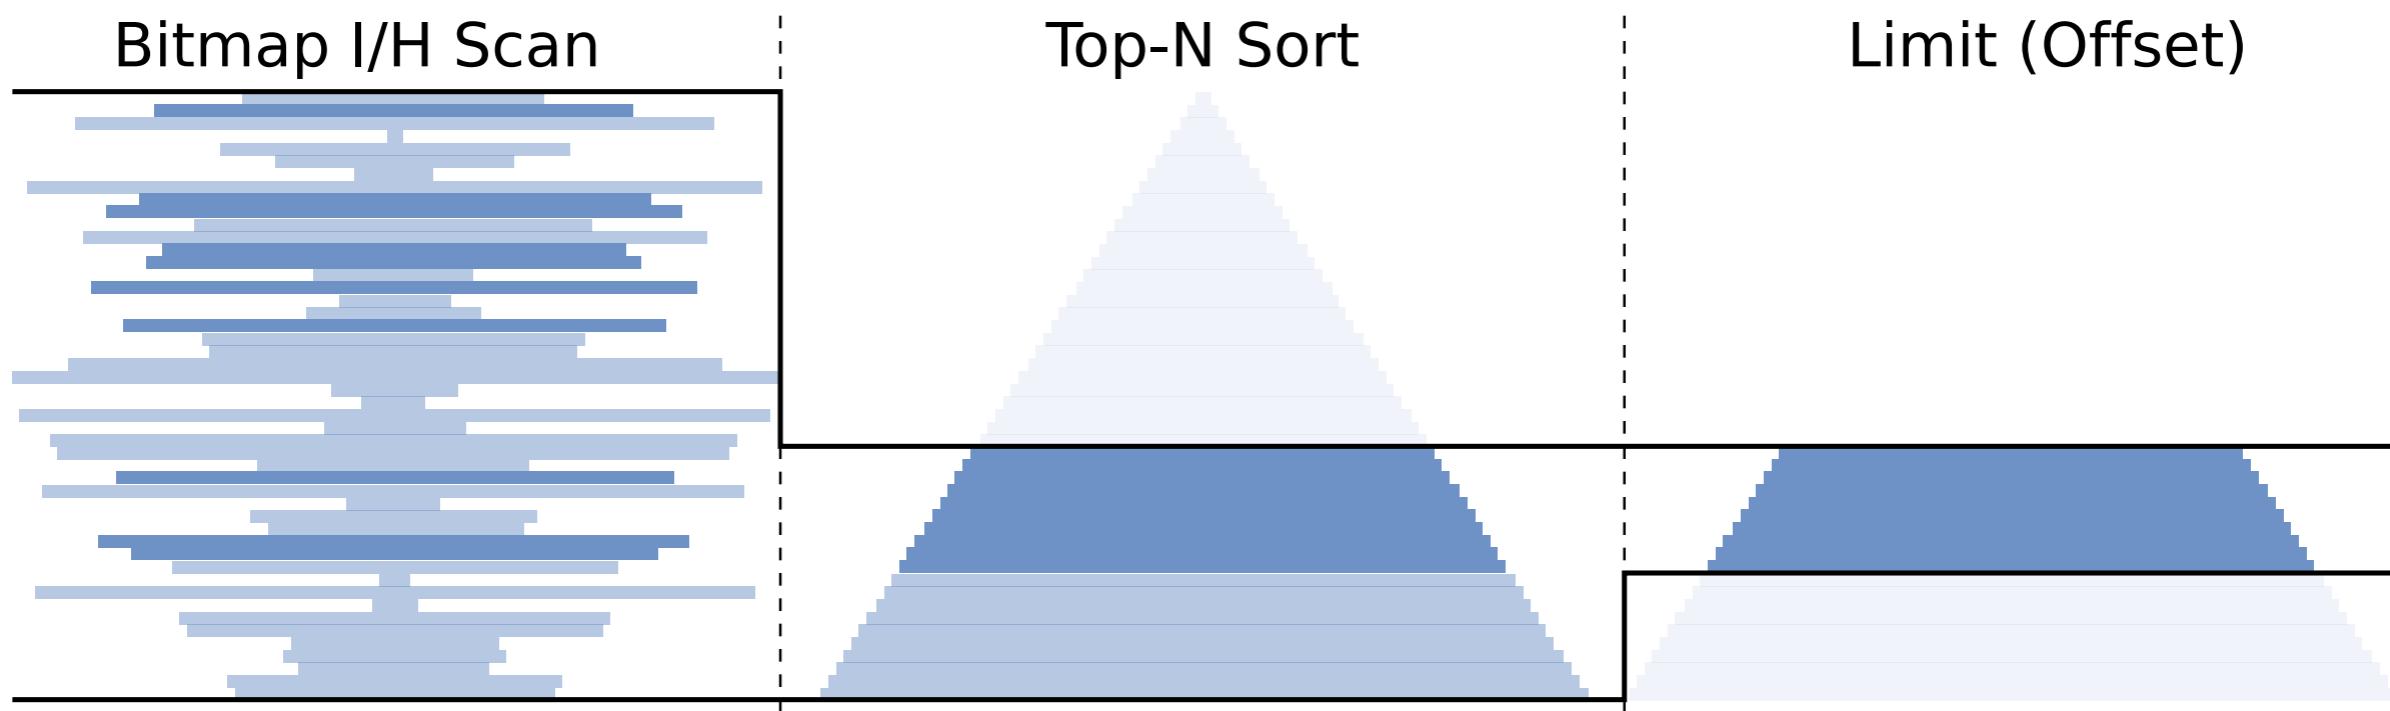


Another Benchmark: Fetch Next Page

Fetching the next page is easy using the offset keyword:

```
select *
  from news
 where topic = 1234
 order by date desc, id desc
offset 10
 limit 10;
```

Worst Case: No Index for order by



Limit (actual rows=10)

-> Sort (**actual rows=20**)

Sort Method: **top-N heapsort Memory: 19kB**

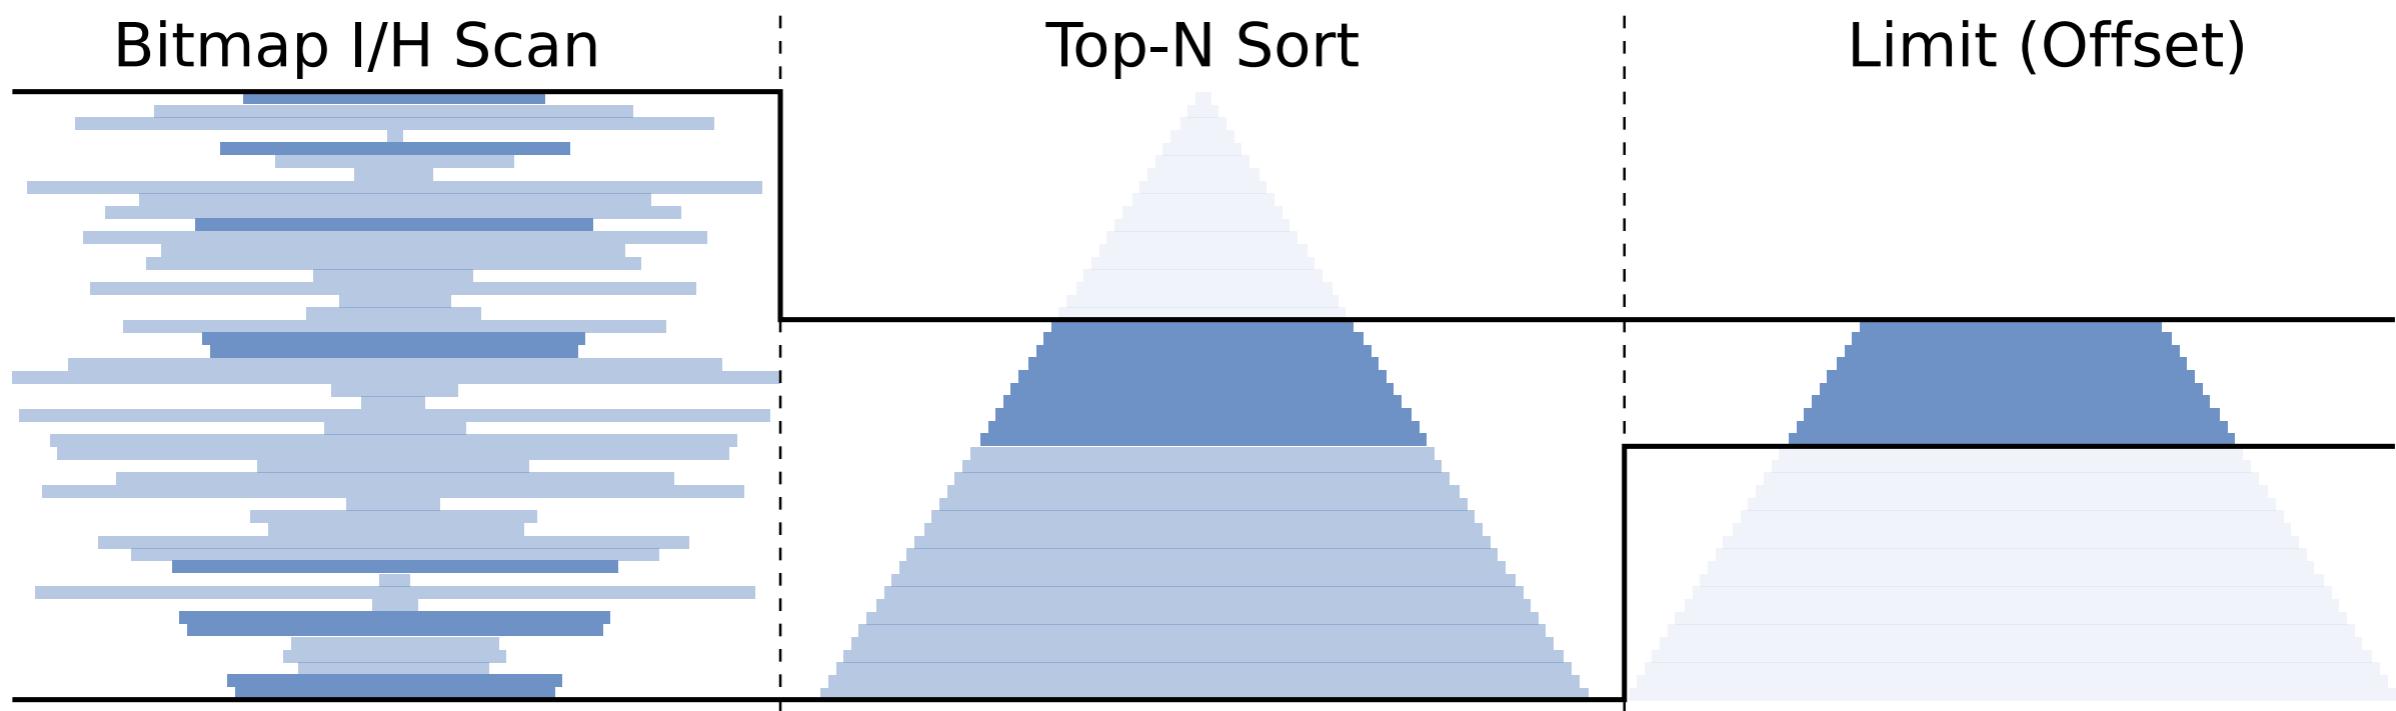
-> Bitmap Heap Scan (actual rows=10000)

Recheck Cond: (topic = 1234)

-> Bitmap Index Scan (actual rows=10000)

Index Cond: (topic = 1234)

Worst Case: No Index for order by



Limit (actual rows=10)

-> Sort (**actual rows=30**)

Sort Method: **top-N heapsort Memory: 20kB**

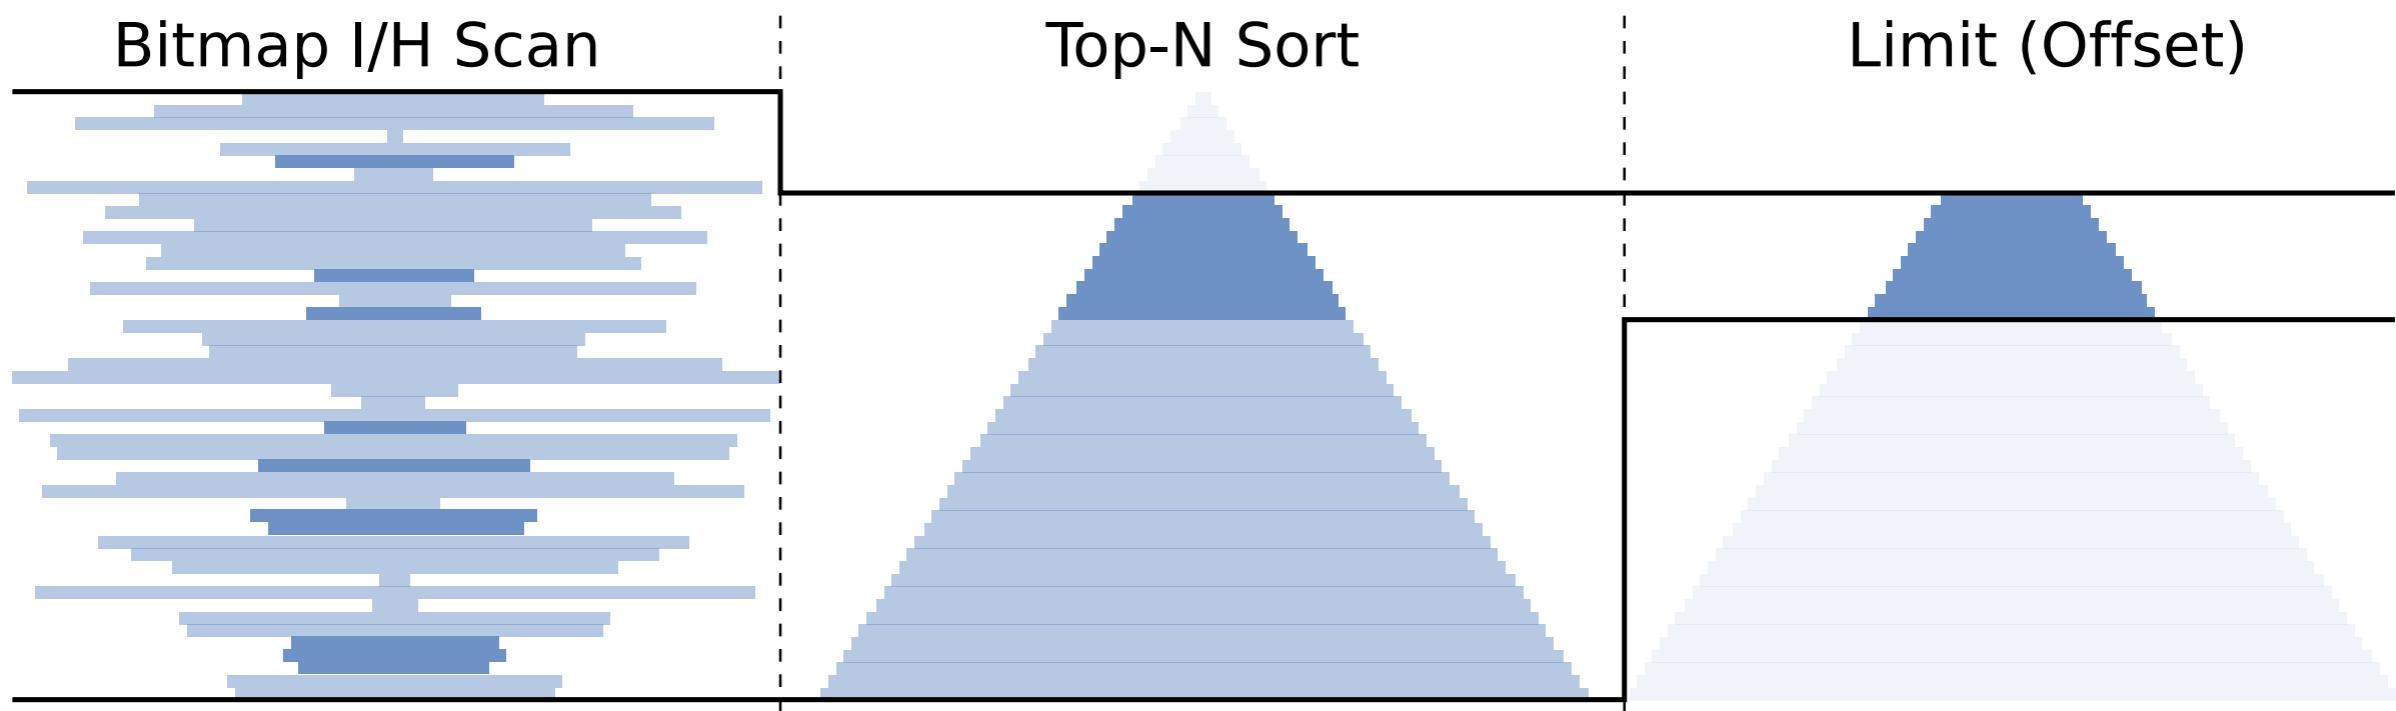
-> Bitmap Heap Scan (actual rows=10000)

Recheck Cond: (topic = 1234)

-> Bitmap Index Scan (actual rows=10000)

Index Cond: (topic = 1234)

Worst Case: No Index for order by



Limit (actual rows=10)

-> Sort (**actual rows=40**)

Sort Method: **top-N heapsort Memory: 22kB**

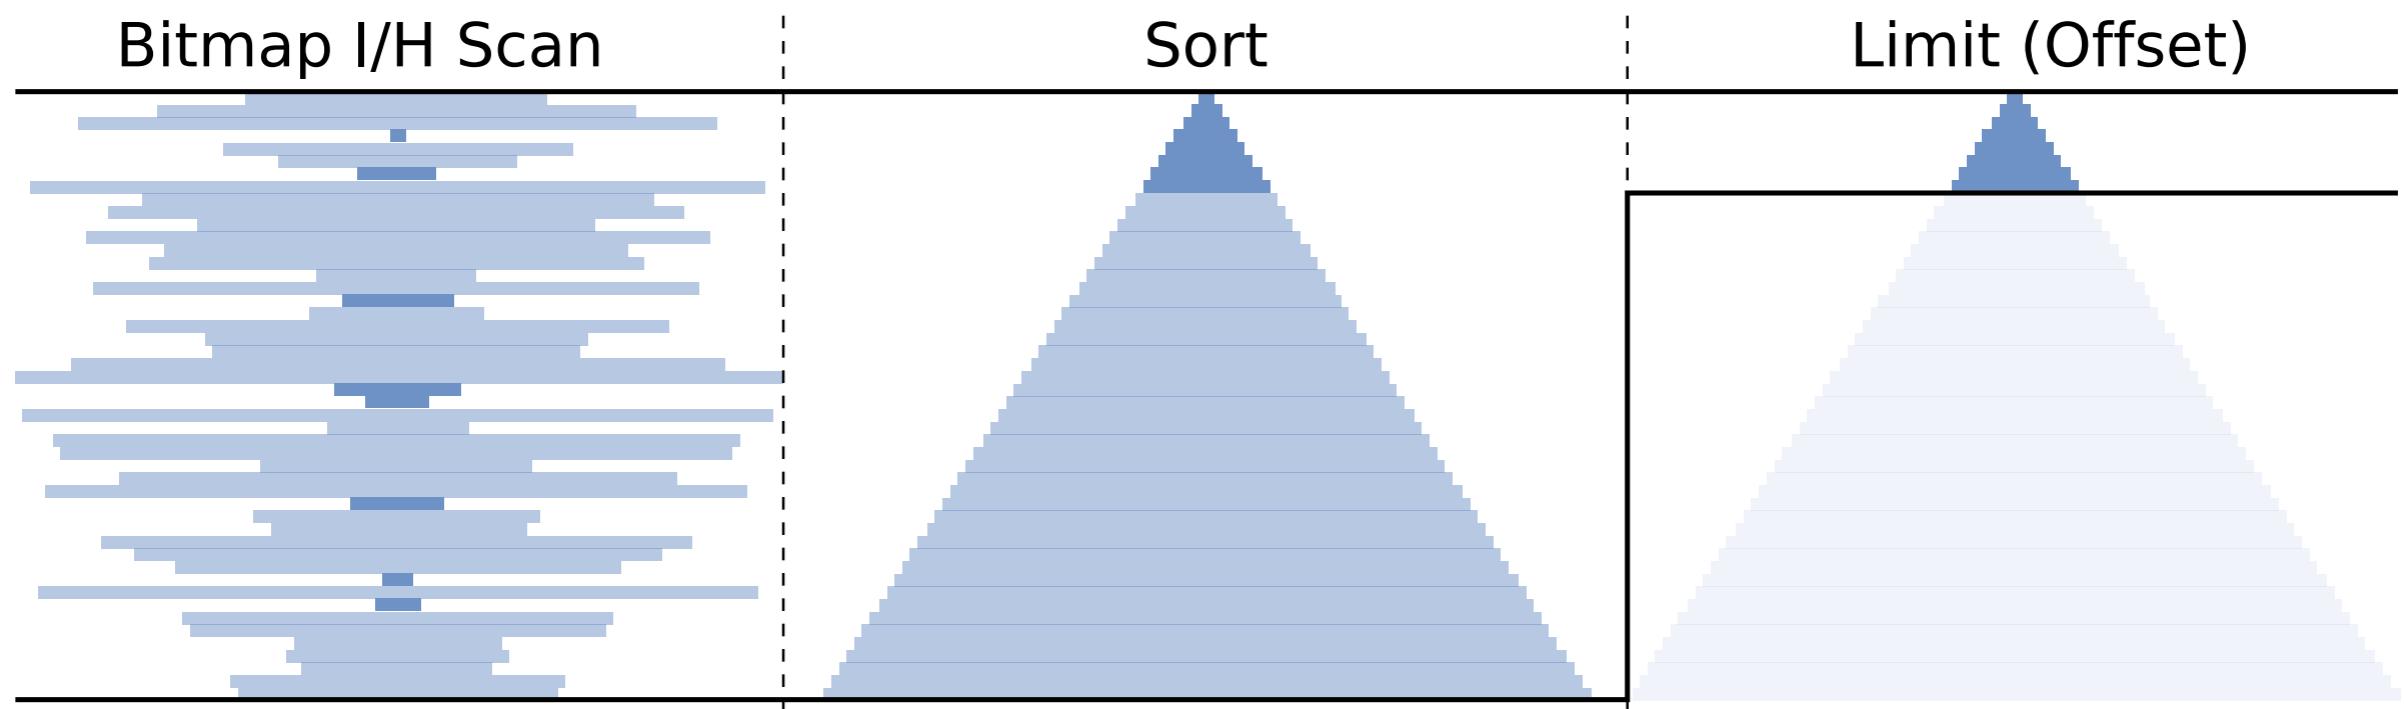
-> Bitmap Heap Scan (actual rows=10000)

Recheck Cond: (topic = 1234)

-> Bitmap Index Scan (actual rows=10000)

Index Cond: (topic = 1234)

Worst Case: No Index for order by



Limit (actual rows=10)

-> Sort (**actual rows=10000**)

Sort Method: **external merge Disk: 1200kB**

-> Bitmap Heap Scan (actual rows=10000)

Recheck Cond: (topic = 1234)

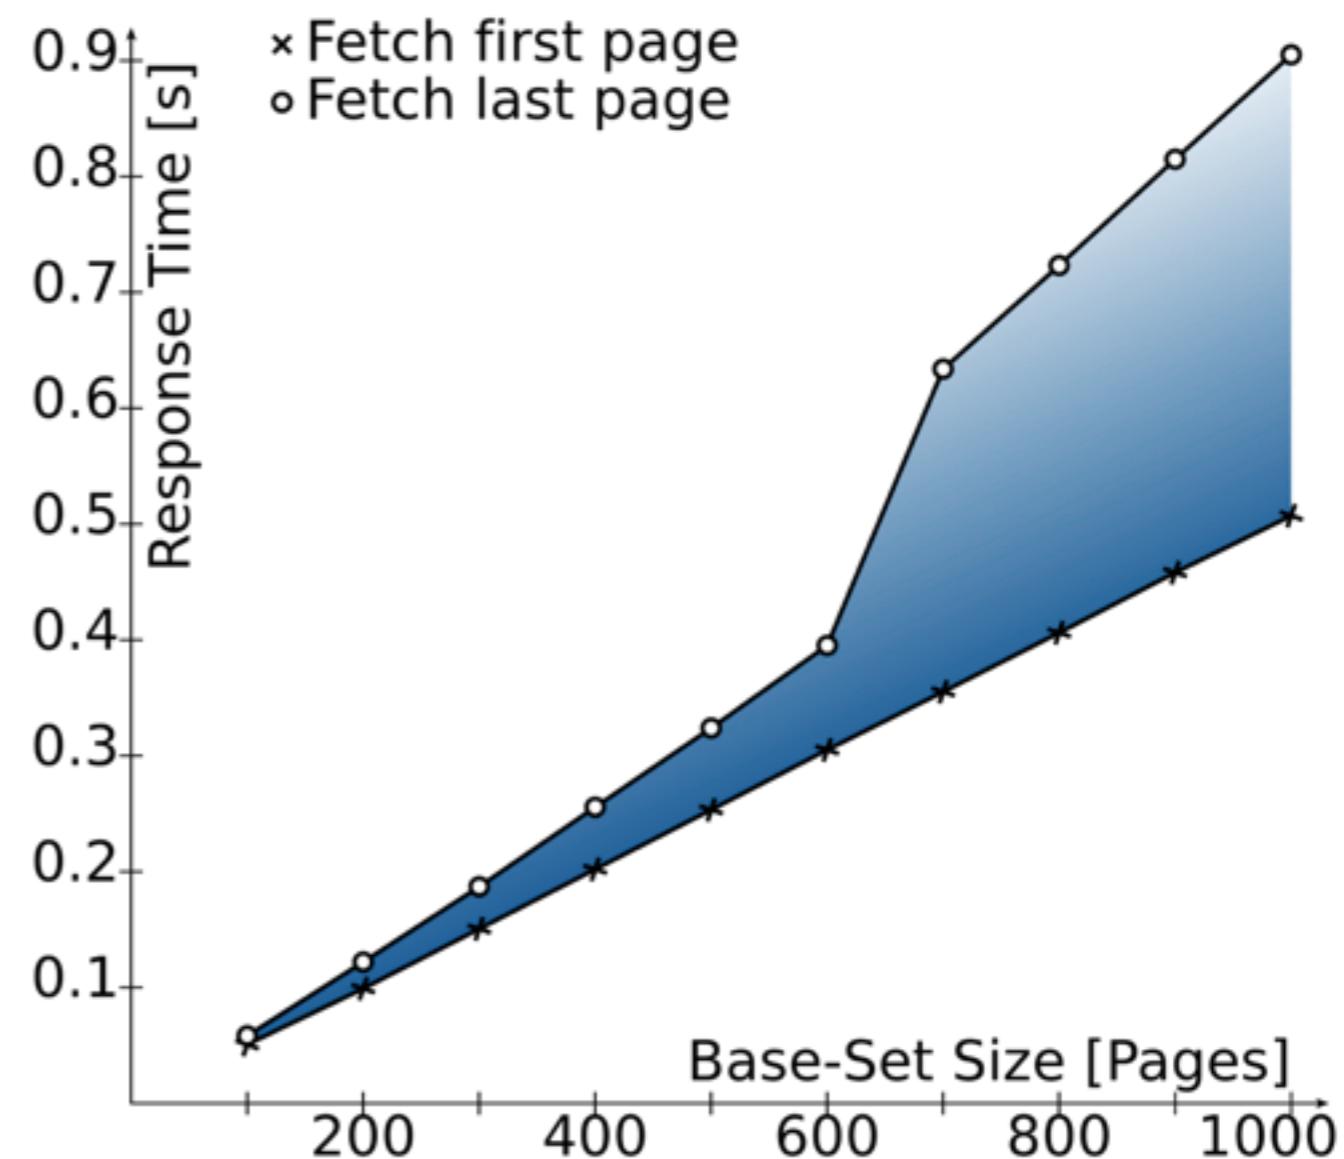
-> Bitmap Index Scan (actual rows=10000)

Index Cond: (topic = 1234)

Worst Case: No Index for order by

Sorting might become the limiting factor when browsing farther back.

Fetching the last page can take considerably longer than fetching the first page.



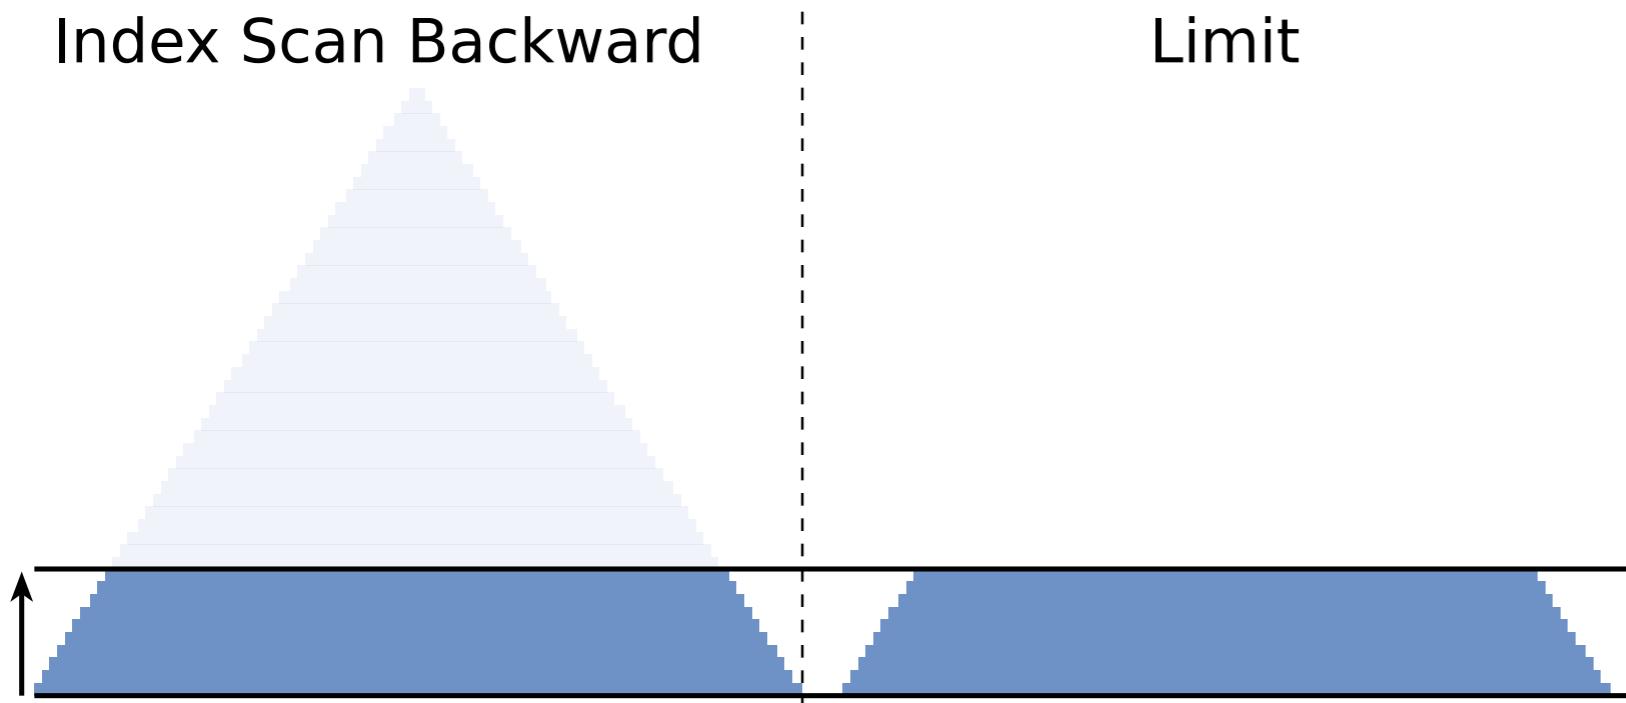
Improvement 1: Indexed order by

```
select *
  from news
 where topic = 1234
 order by date desc, id desc
offset 10
 limit 10;
```

create index .. on news (topic, date, id);

A single index to support the where and order by clauses.

Improvement 1: Indexed order by

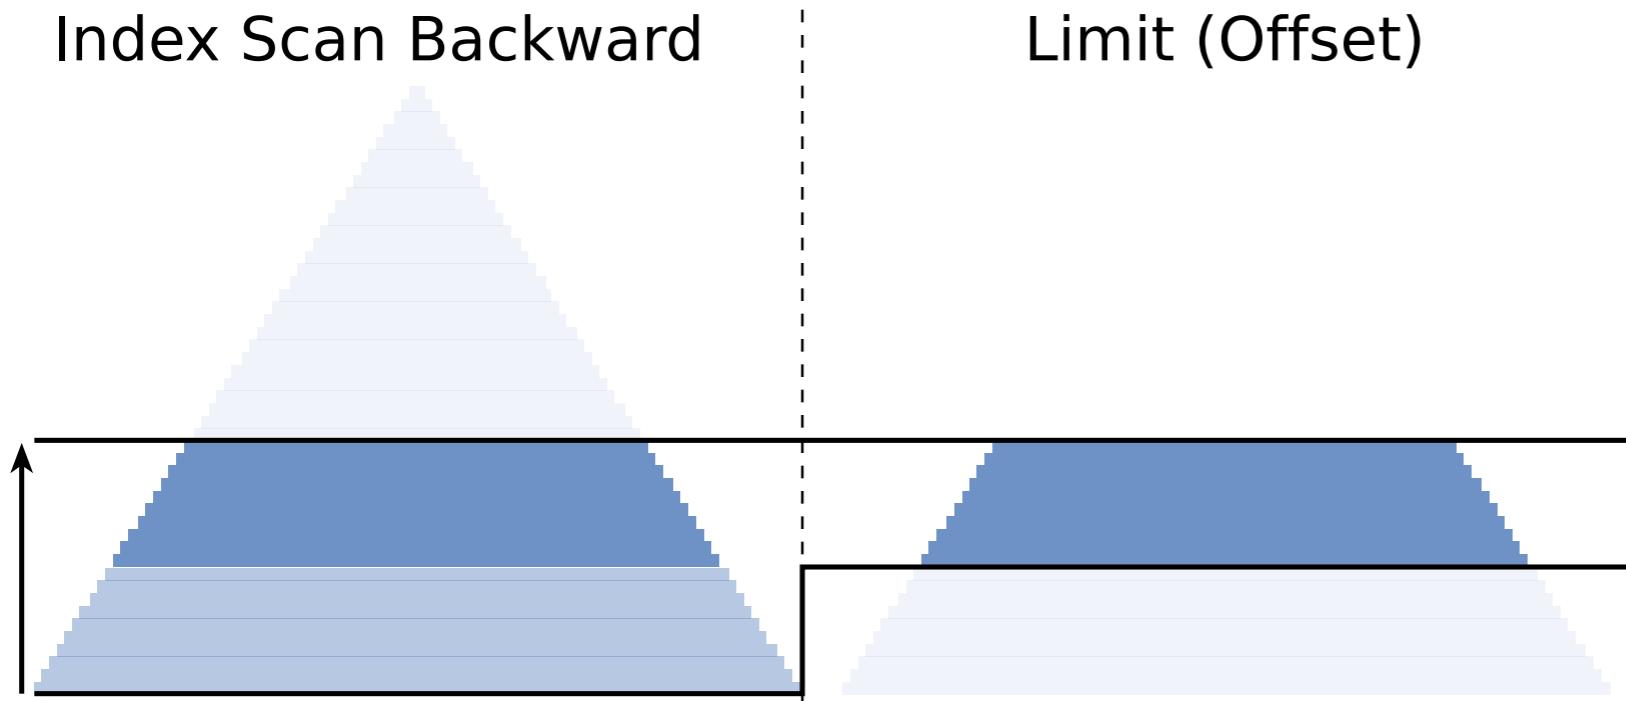


Limit (**actual rows=10**)

-> Index Scan Backward (**actual rows=10**)

Index Cond: (topic = 0)

Improvement 1: Indexed order by

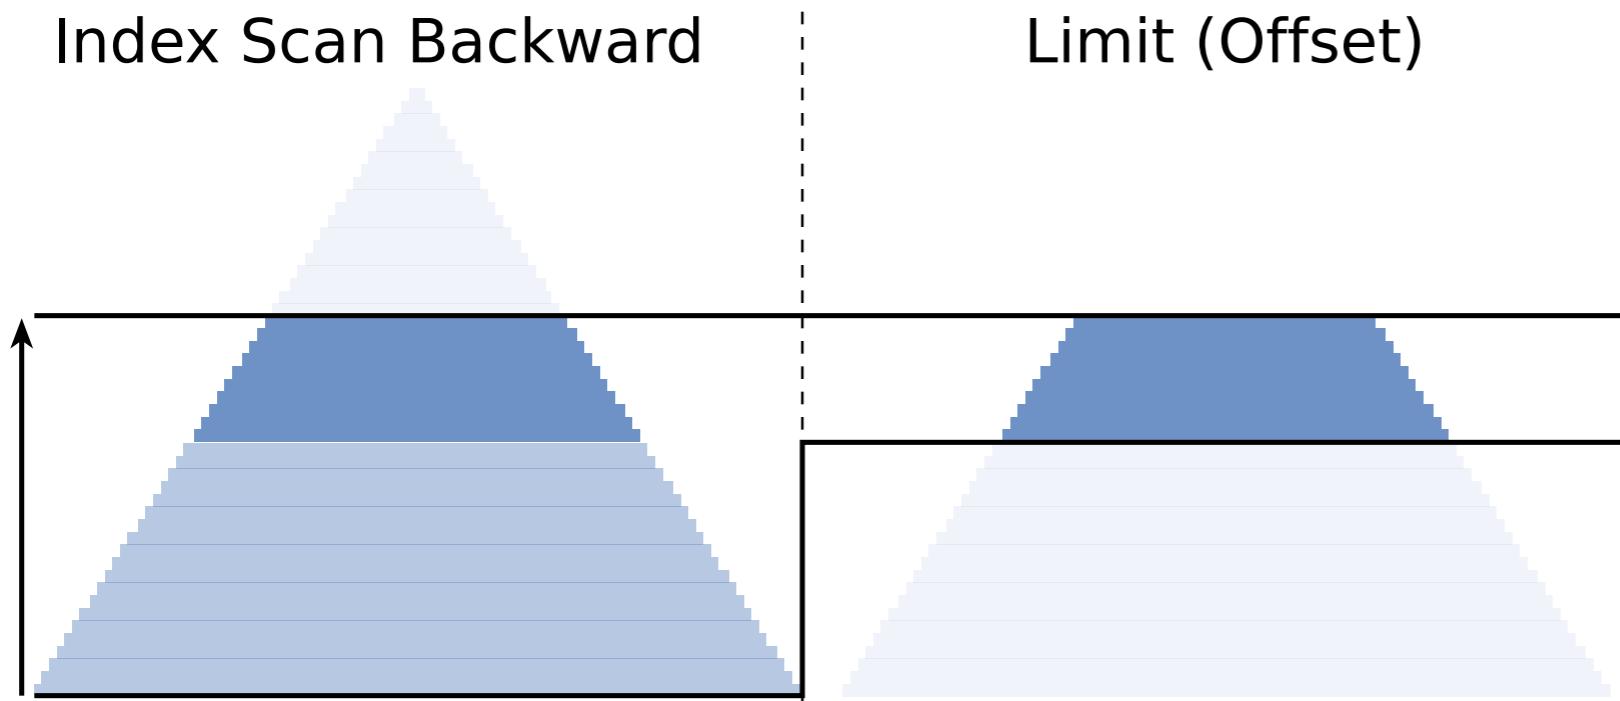


Limit (**actual rows=10**)

-> Index Scan Backward (**actual rows=20**)

Index Cond: (topic = 0)

Improvement 1: Indexed order by

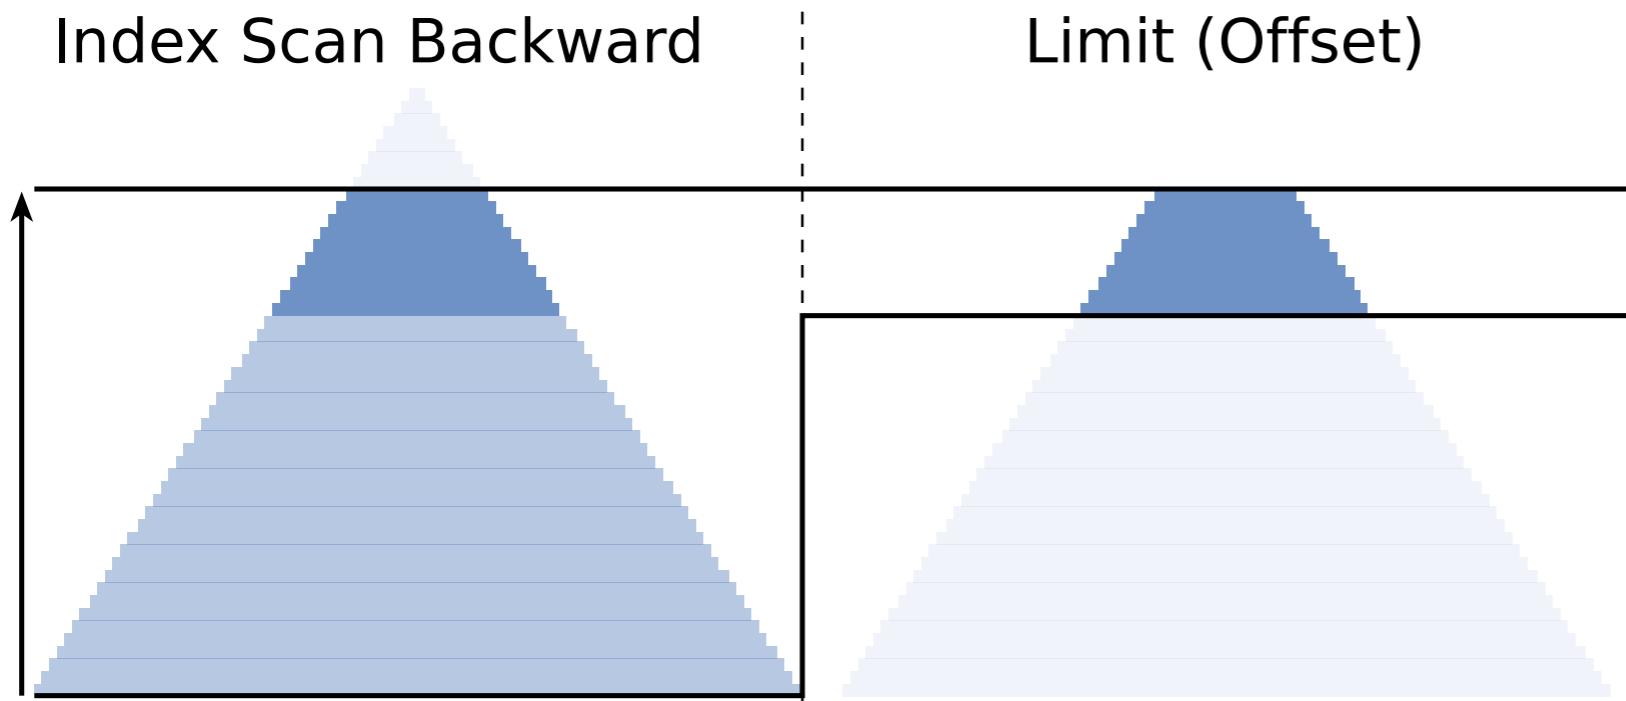


Limit (**actual rows=10**)

-> Index Scan Backward (**actual rows=30**)

Index Cond: (topic = 0)

Improvement 1: Indexed order by

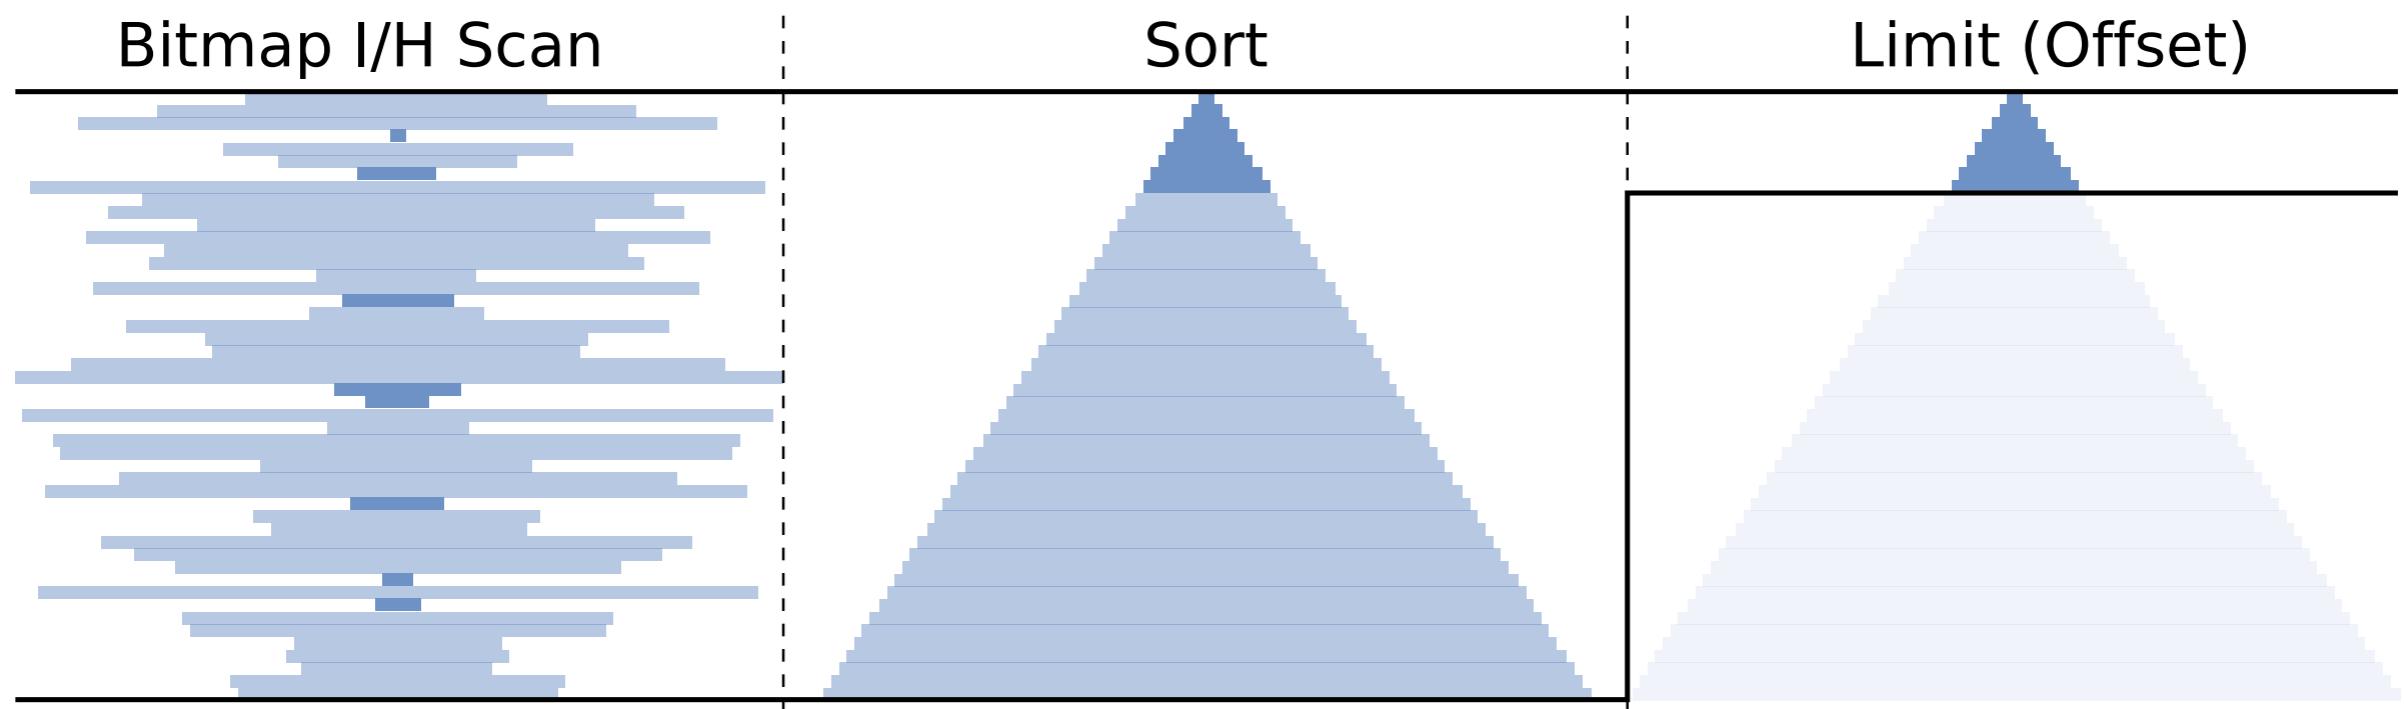


Limit (**actual rows=10**)

-> Index Scan Backward (**actual rows=40**)

Index Cond: (topic = 0)

Improvement 1: Indexed order by



Limit (actual rows=10)

-> Sort (**actual rows=10000**)

Sort Method: **external merge Disk: 1200kB**

-> Bitmap Heap Scan (actual rows=10000)

Recheck Cond: (topic = 1234)

-> Bitmap Index Scan (actual rows=10000)

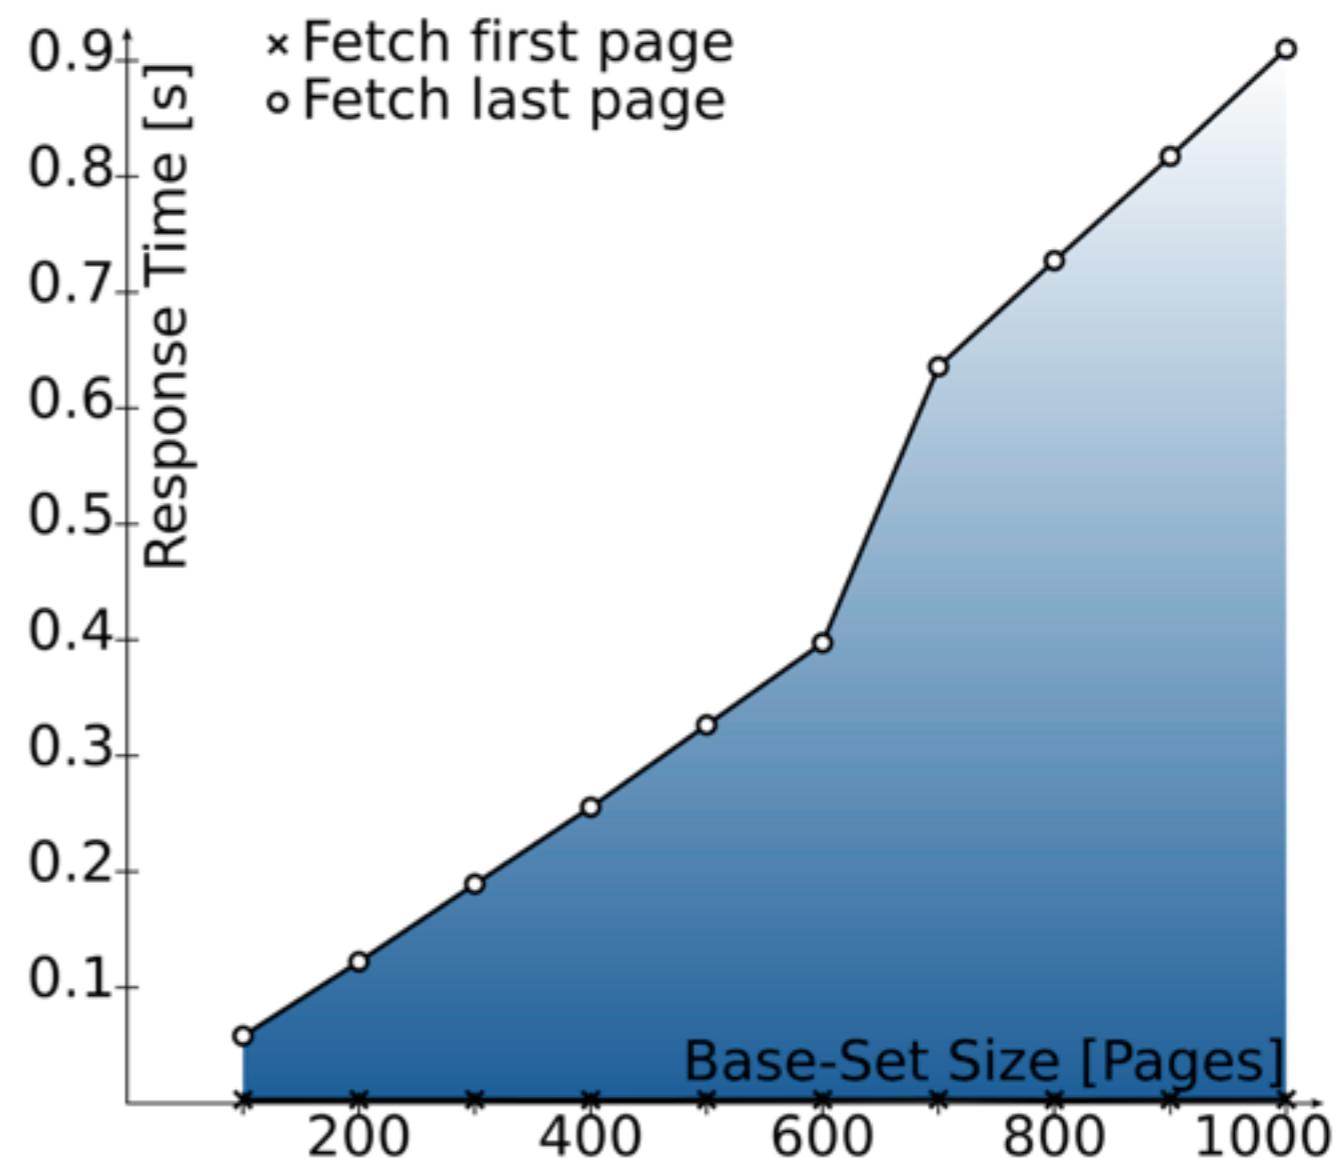
Index Cond: (topic = 1234)

Improvement 1: Indexed order by

Fetching the first page is not affected by the Base-Set size!

Fetching the next page is also faster.

However, PostgreSQL might take a Bitmap Index Scan when browsing to the end.



We can do better!

Don't touch what you don't need

Improvement 2: The Seek Method

Instead of offset, use a where filter to remove the rows from previous pages.

```
select *
  from news
 where topic = 1234
   and (date, id) < (prev_date, prev_id)
 order by date desc, id desc
 limit 10;
```

Only select the rows “before” (=earlier date) the last row from the previous page.

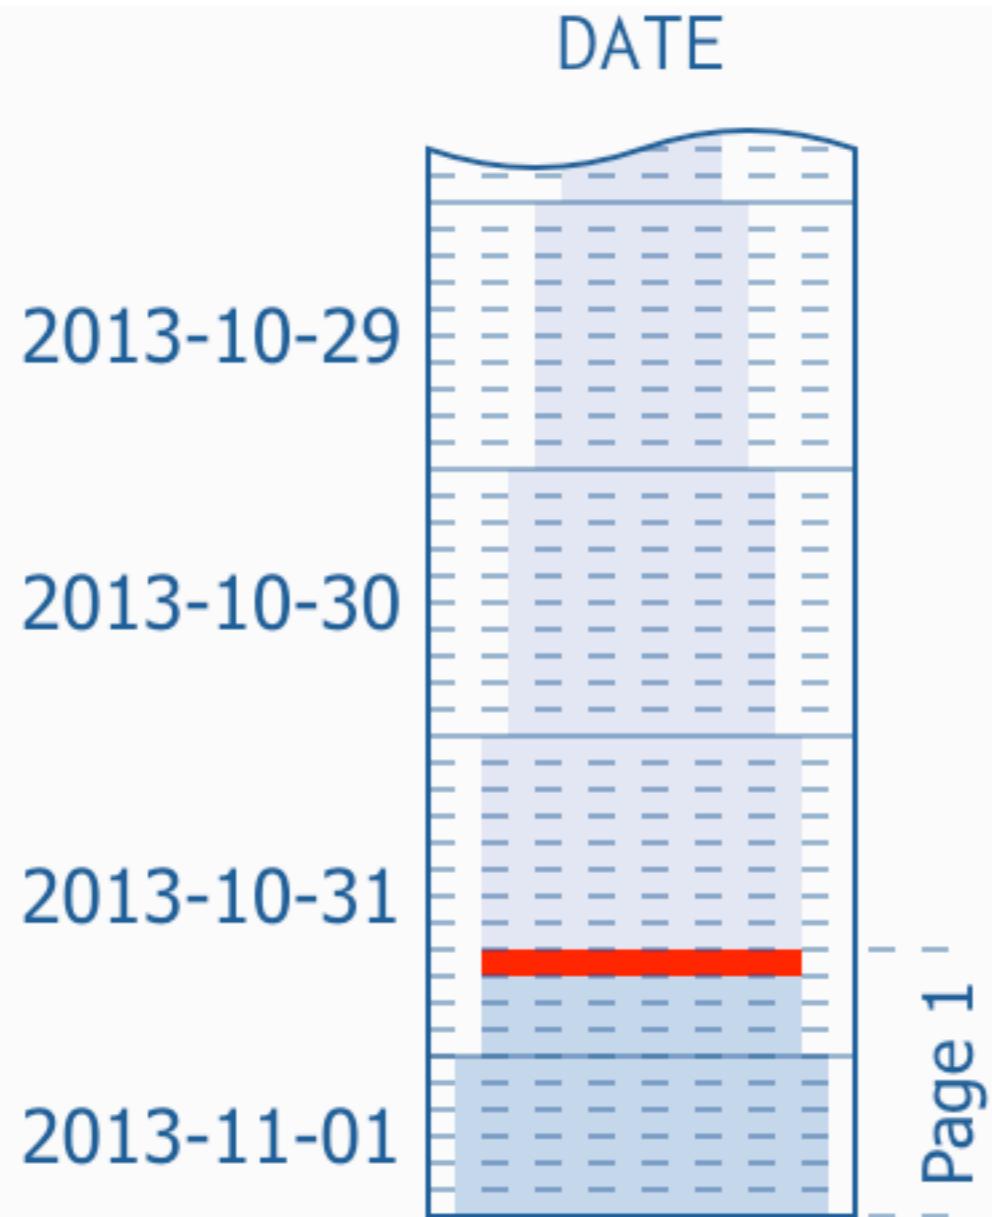
A definite sort order is really required!

Side Note: Row Values

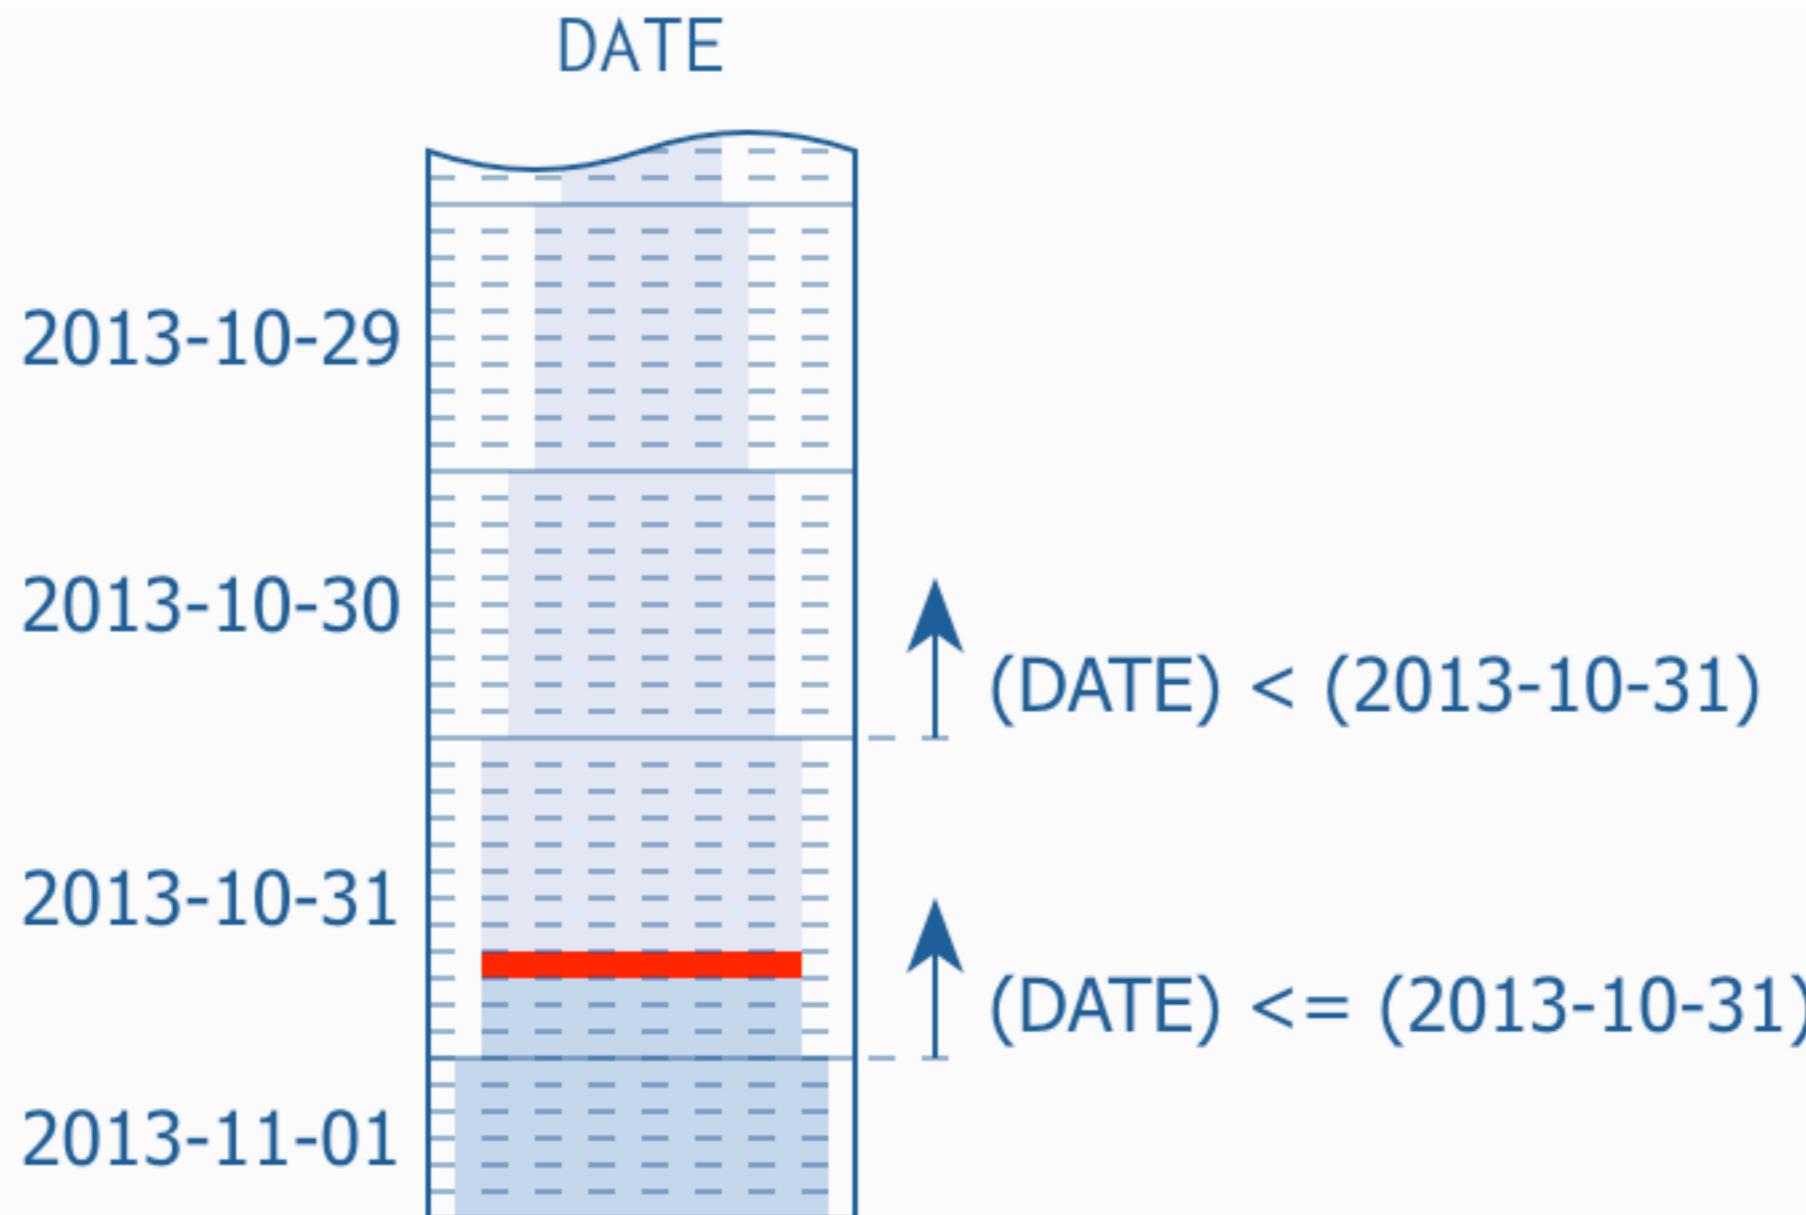
Besides scalar values, SQL also defines “row values” or “composite values.”

- ▶ In the SQL standard since ages (SQL-92)
- ▶ All comparison operators are well defined
 - ▶ E.g.: $(x, y) > (a, b)$ is true iff
$$(x > a \text{ or } (x=a \text{ and } y>b))$$
 - ▶ In other words, when (x,y) sorts after (a,b)
- ▶ Great PostgreSQL support since 8.0!

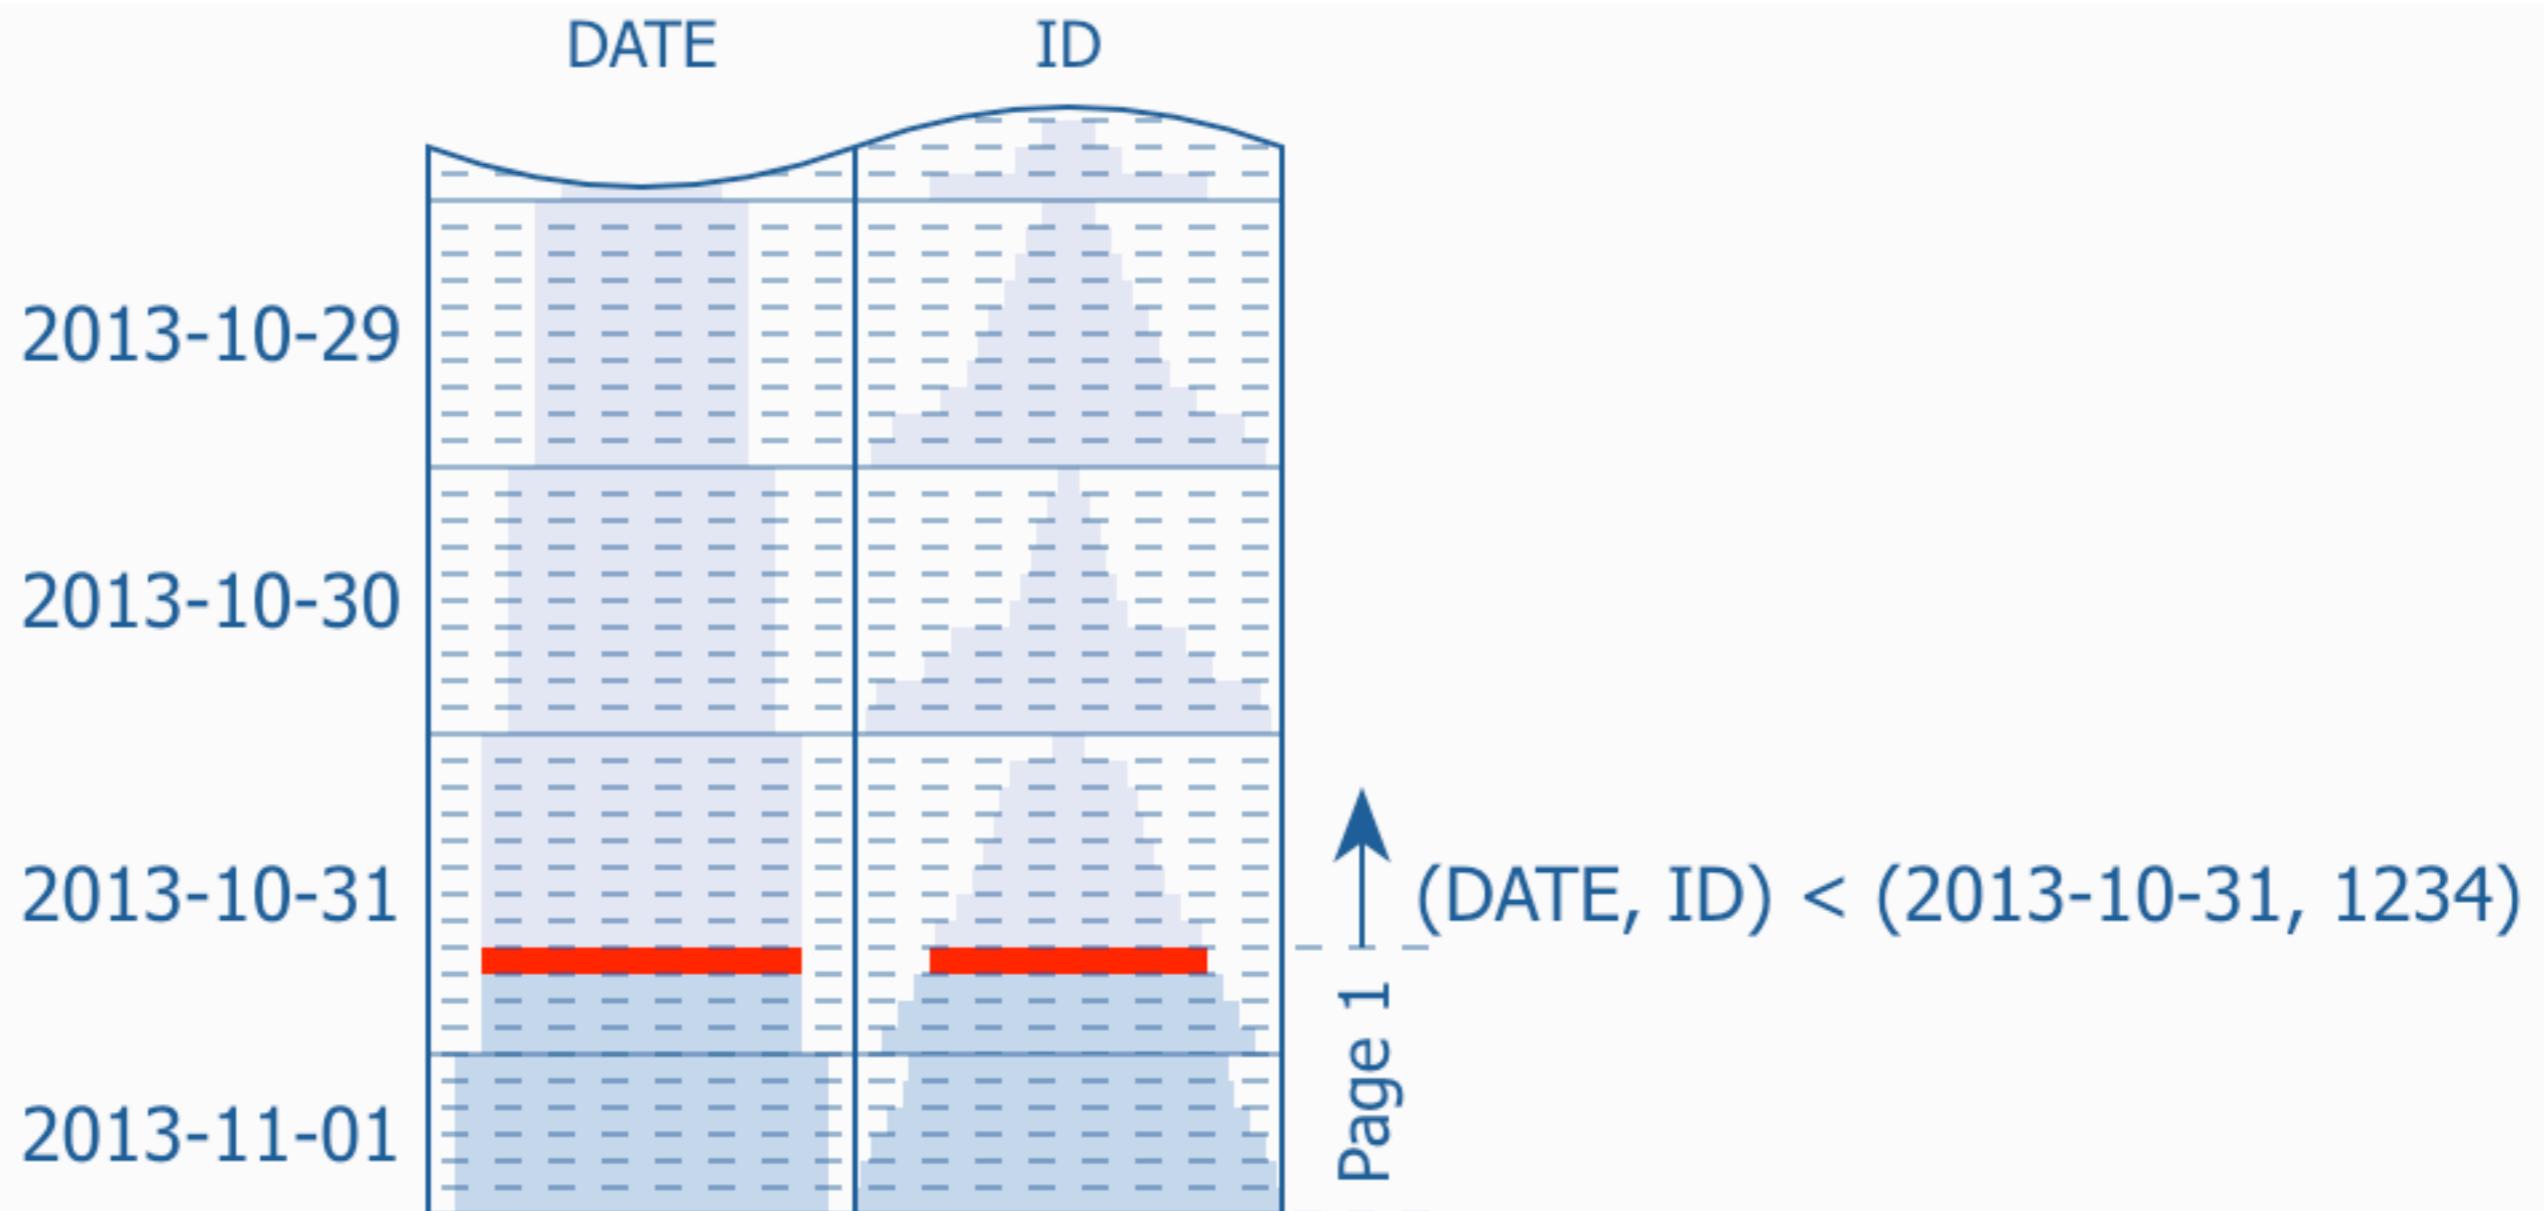
Row Values: How it works



Row Values: How it works



Row Values: How it works



Row Values: Emulating them

- ▶ There is only a single database that has sufficient Row-Values support for this trick (PostgreSQL).
- ▶ With all other database, you must emulate it according to this scheme:

```
WHERE sale_date <= ?
  AND (
    (sale_date < ?)
    OR
    (sale_date = ? AND sale_id < ?)
  )
```

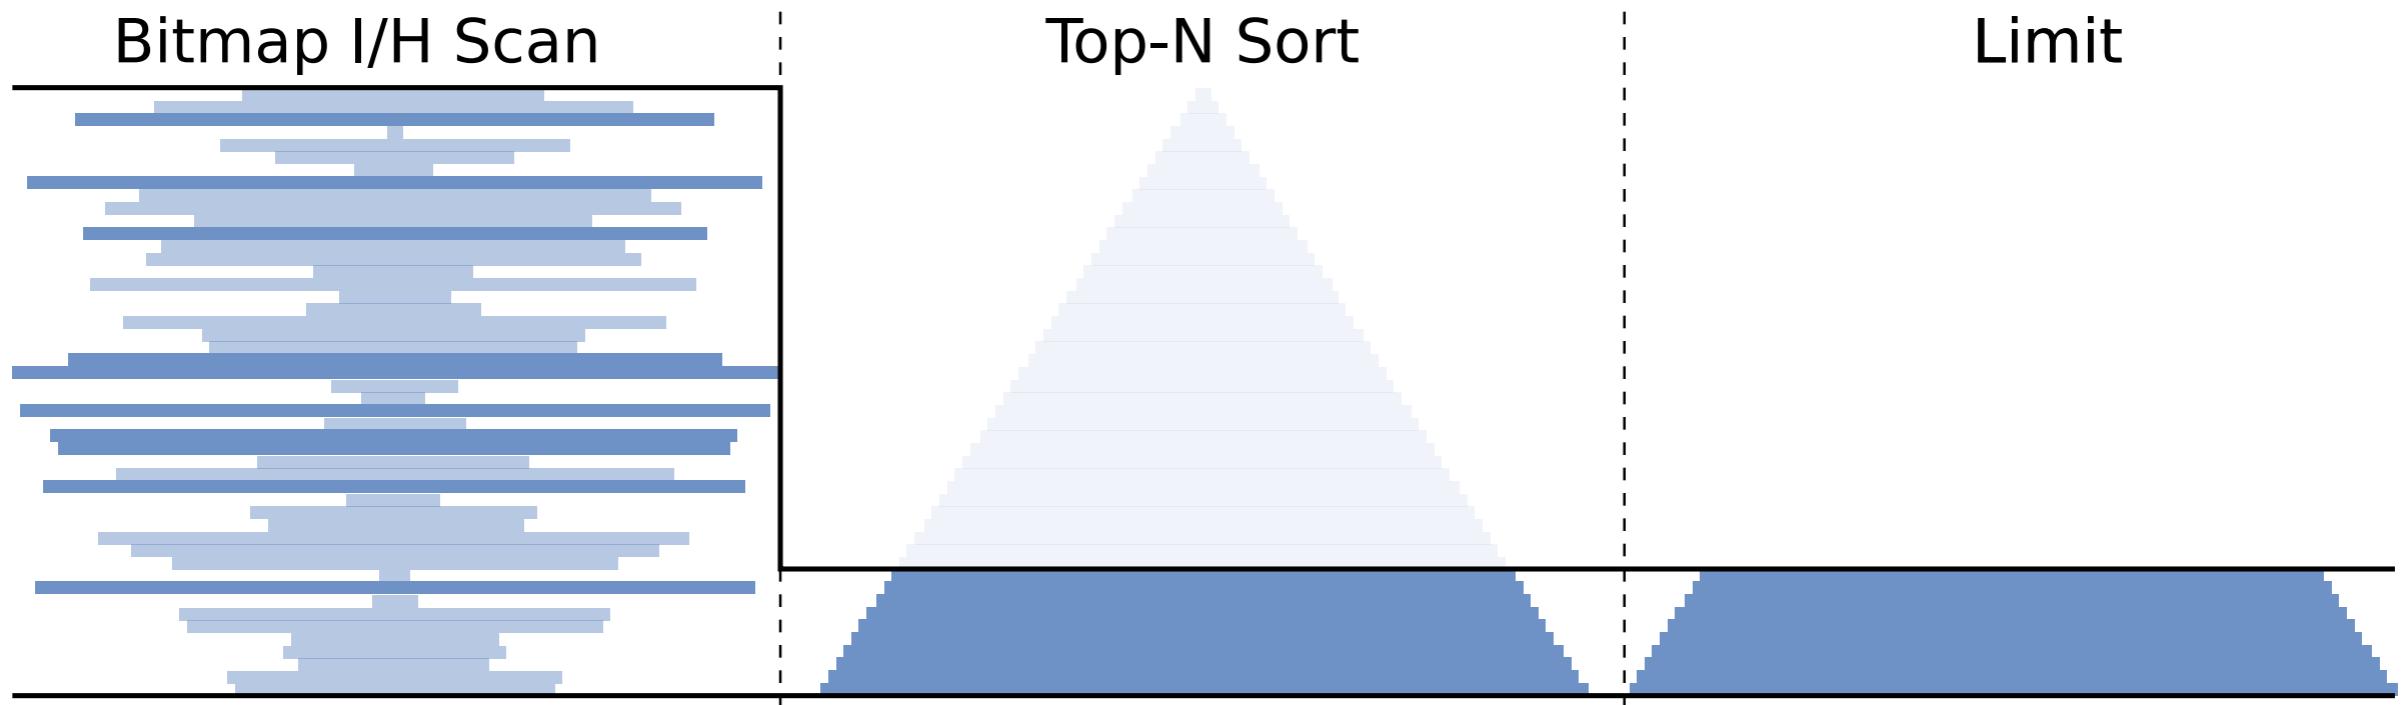
See: <http://use-the-index-luke.com/sql/partial-results/fetch-next-page#sb-equivalent-logic>

Seek-Method without Optimal Index

```
select *
  from news
 where topic = 1234
   and (date, id) < (prev_date, prev_id)
 order by date desc, id desc
 limit 10;
```

create index .. on news (topic);

Seek Method w/o Index for order by



Limit (**actual rows=10**)

-> Sort (**actual rows=10**)

Sort Method: **top-N heapsort** Memory: 18kB

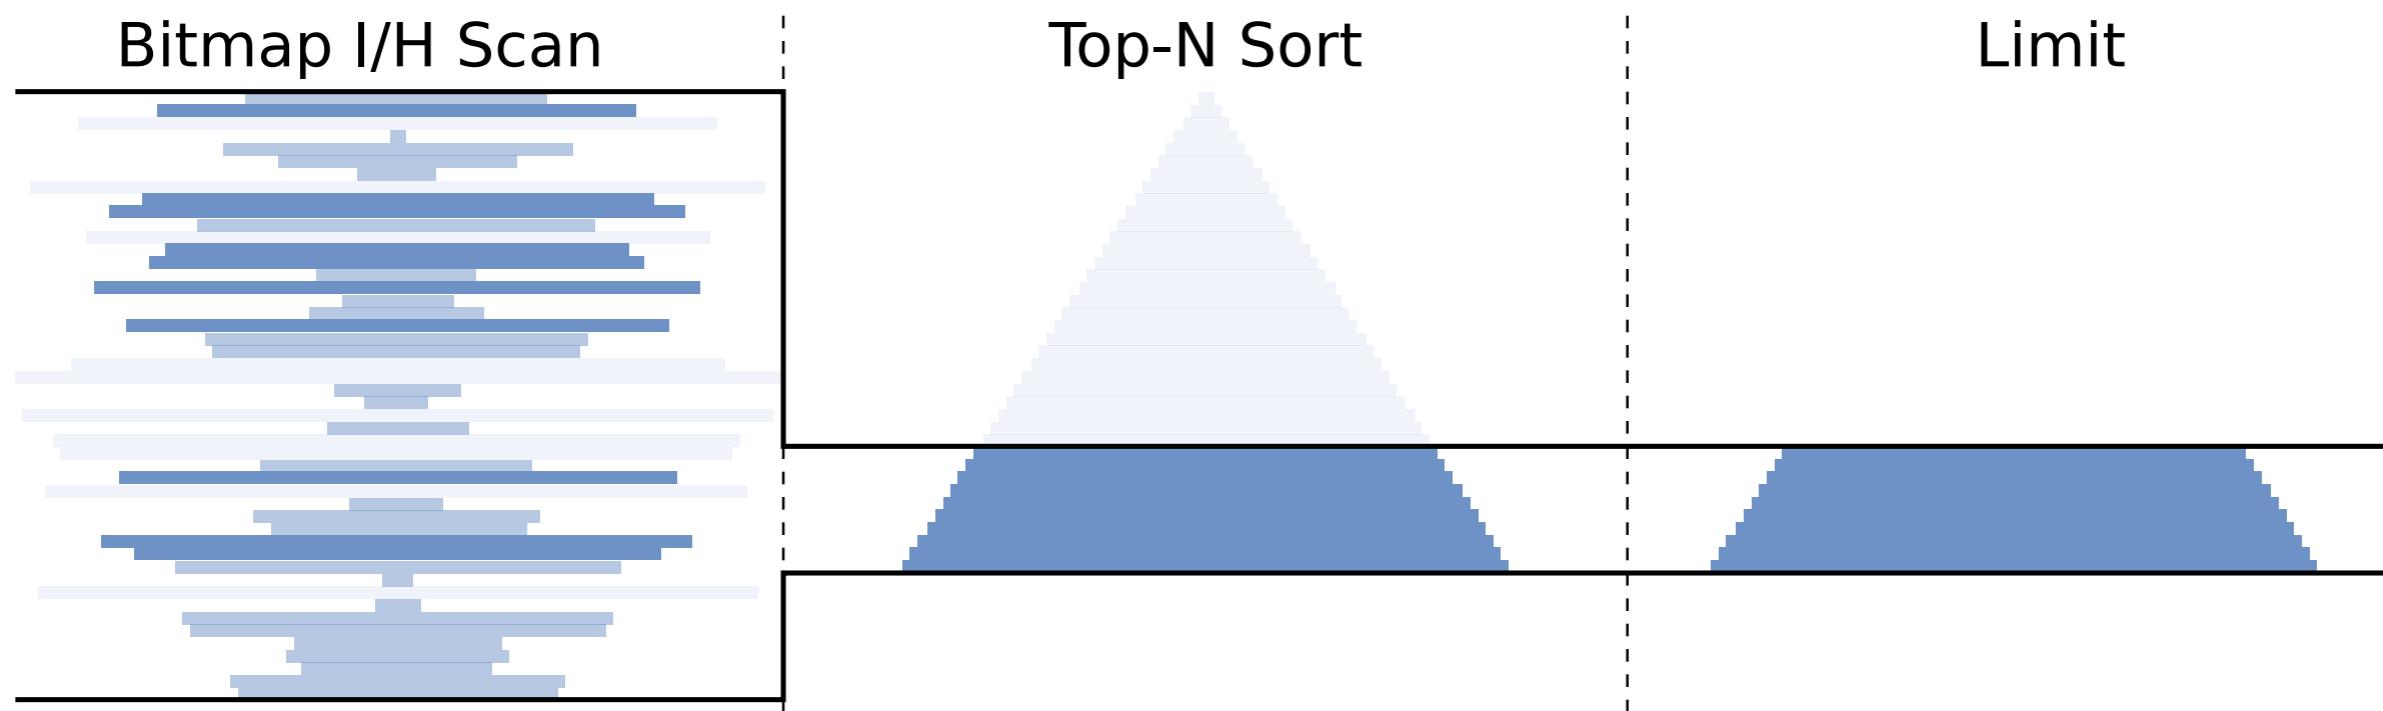
-> Bitmap Heap Scan (**rows=10000**)

Recheck Cond: (topic = 1234)

-> Bitmap Index Scan (**rows=10000**)

Index Cond: (topic = 1234)

Seek Method w/o Index for order by



Limit (actual rows=10)

-> Sort (**actual rows=10**)

Sort Method: **top-N heapsort** Memory: 18kB

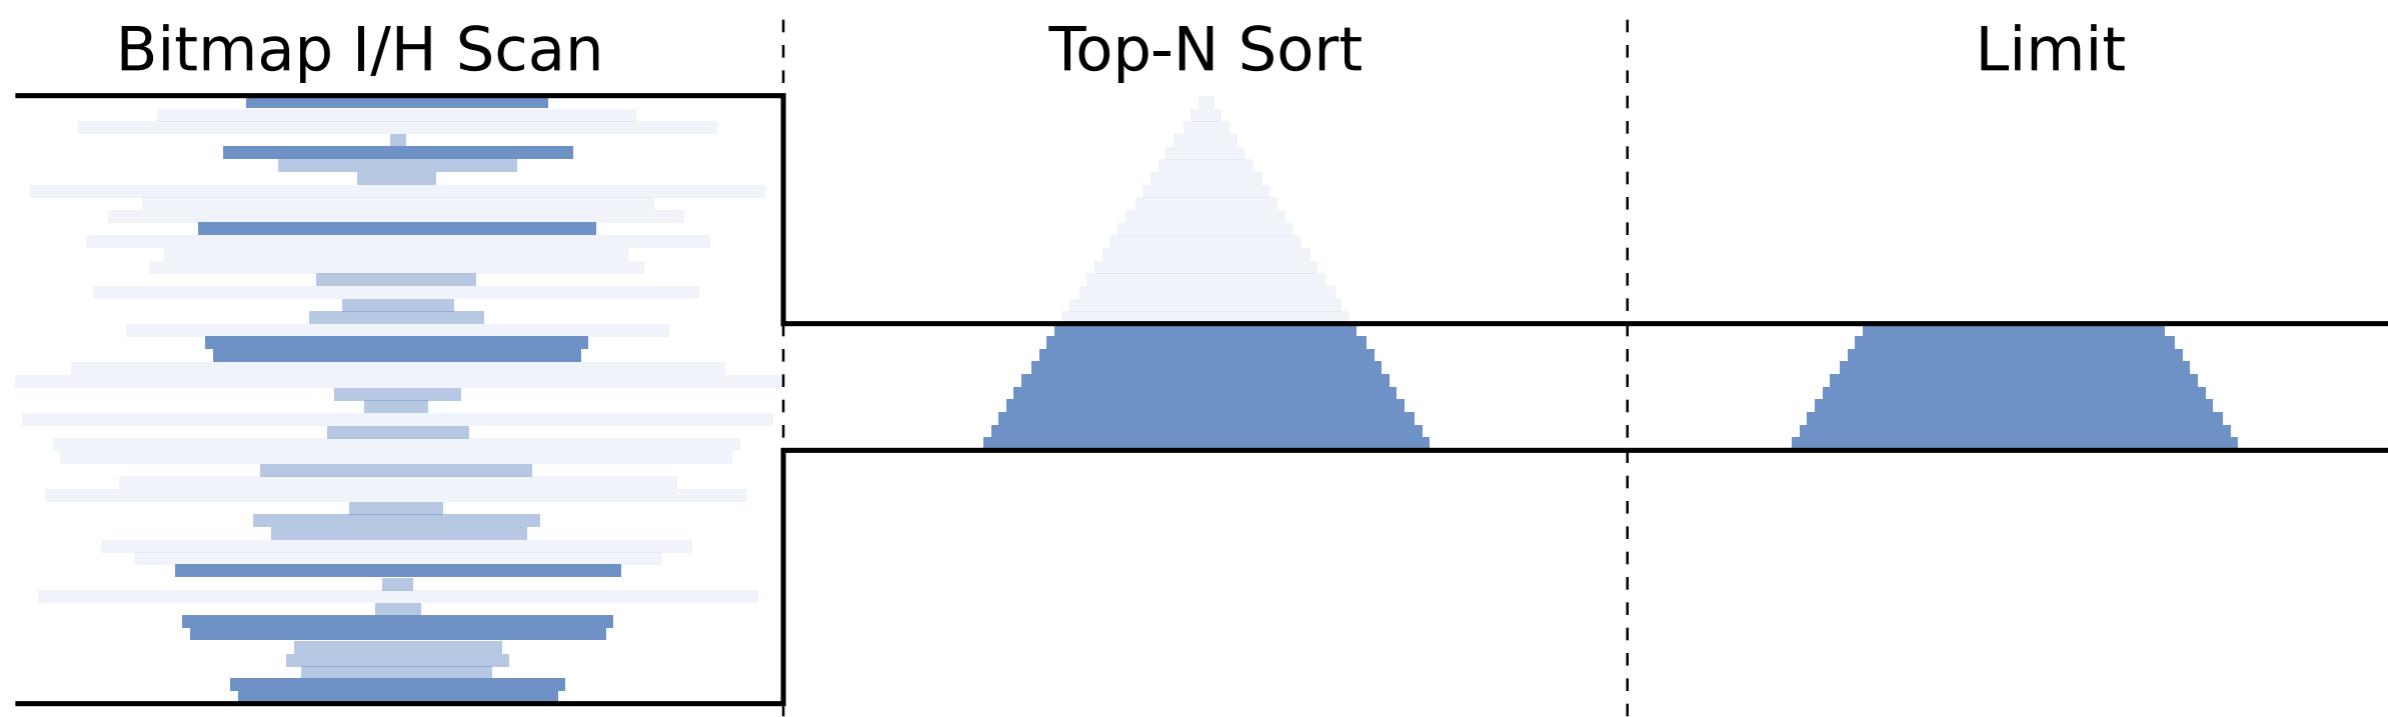
-> Bitmap Heap Scan (**actual rows=9990**)

Rows Removed by Filter: 10 (new in 9.2)

-> Bitmap Index Scan (**actual rows=10000**)

Index Cond: (topic = 1234)

Seek Method w/o Index for order by



Limit (actual rows=10)

-> Sort (**actual rows=10**)

Sort Method: **top-N heapsort** Memory: 18kB

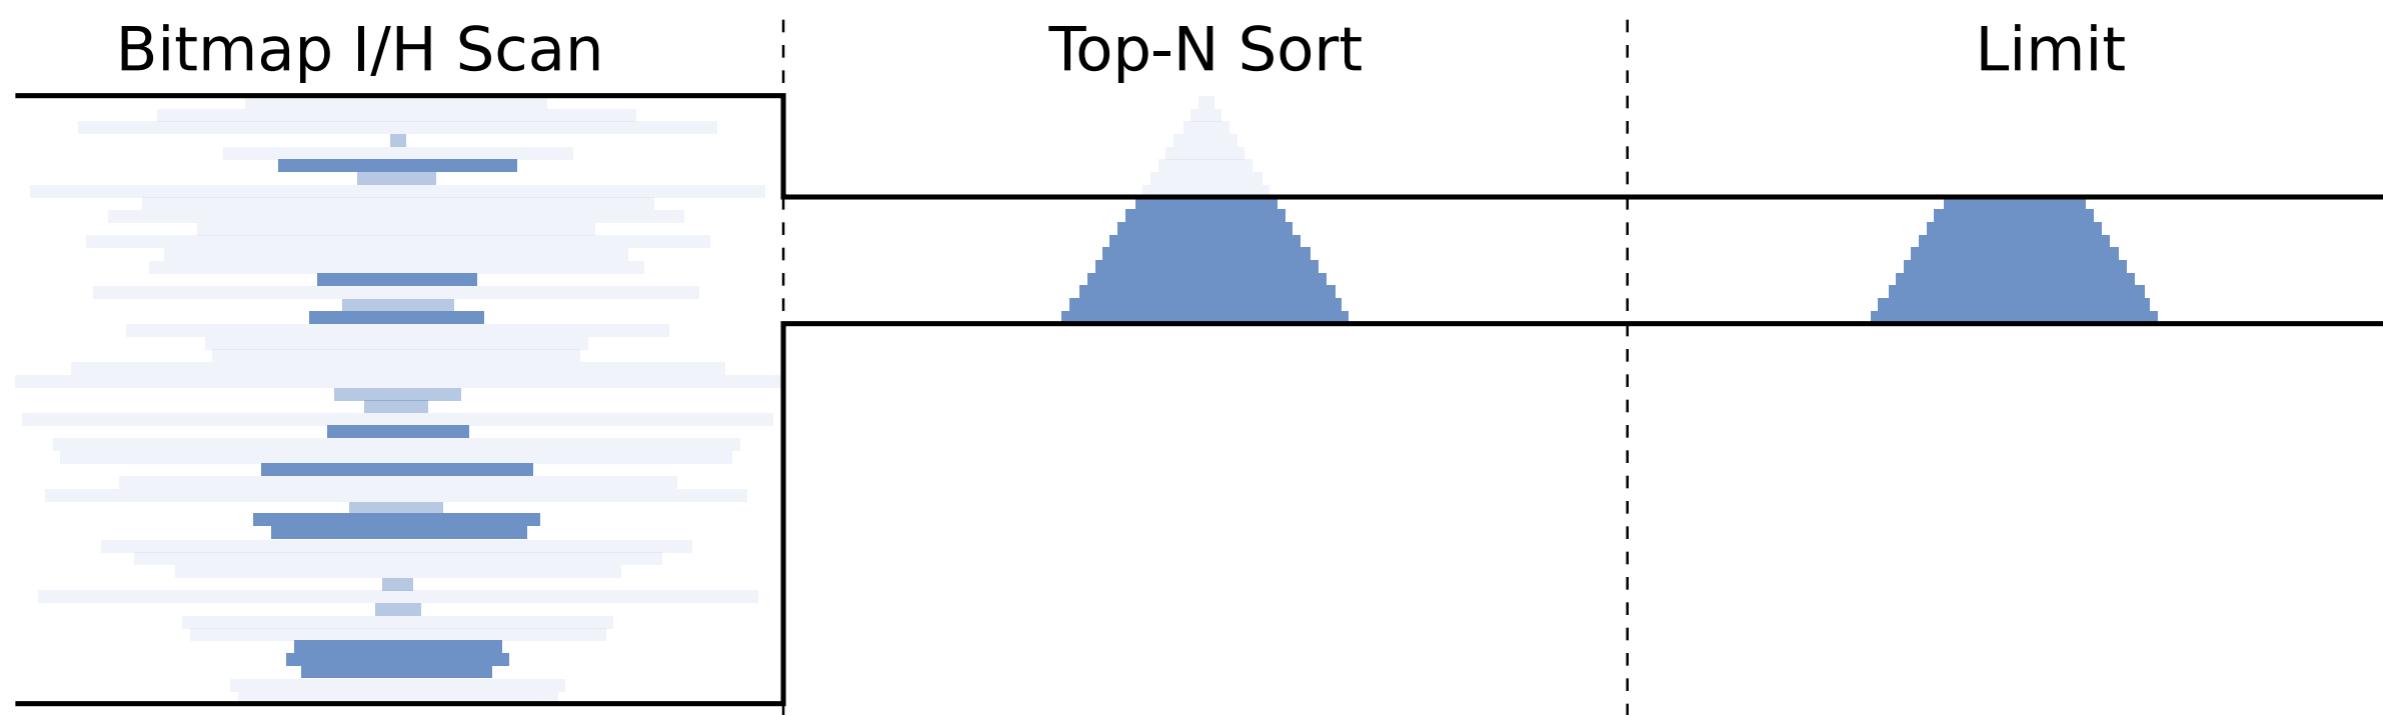
-> Bitmap Heap Scan (**actual rows=9980**)

Rows Removed by Filter: 20 (new in 9.2)

-> Bitmap Index Scan (**actual rows=10000**)

Index Cond: (topic = 1234)

Seek Method w/o Index for order by



Limit (actual rows=10)

-> Sort (**actual rows=10**)

Sort Method: **top-N heapsort** Memory: 18kB

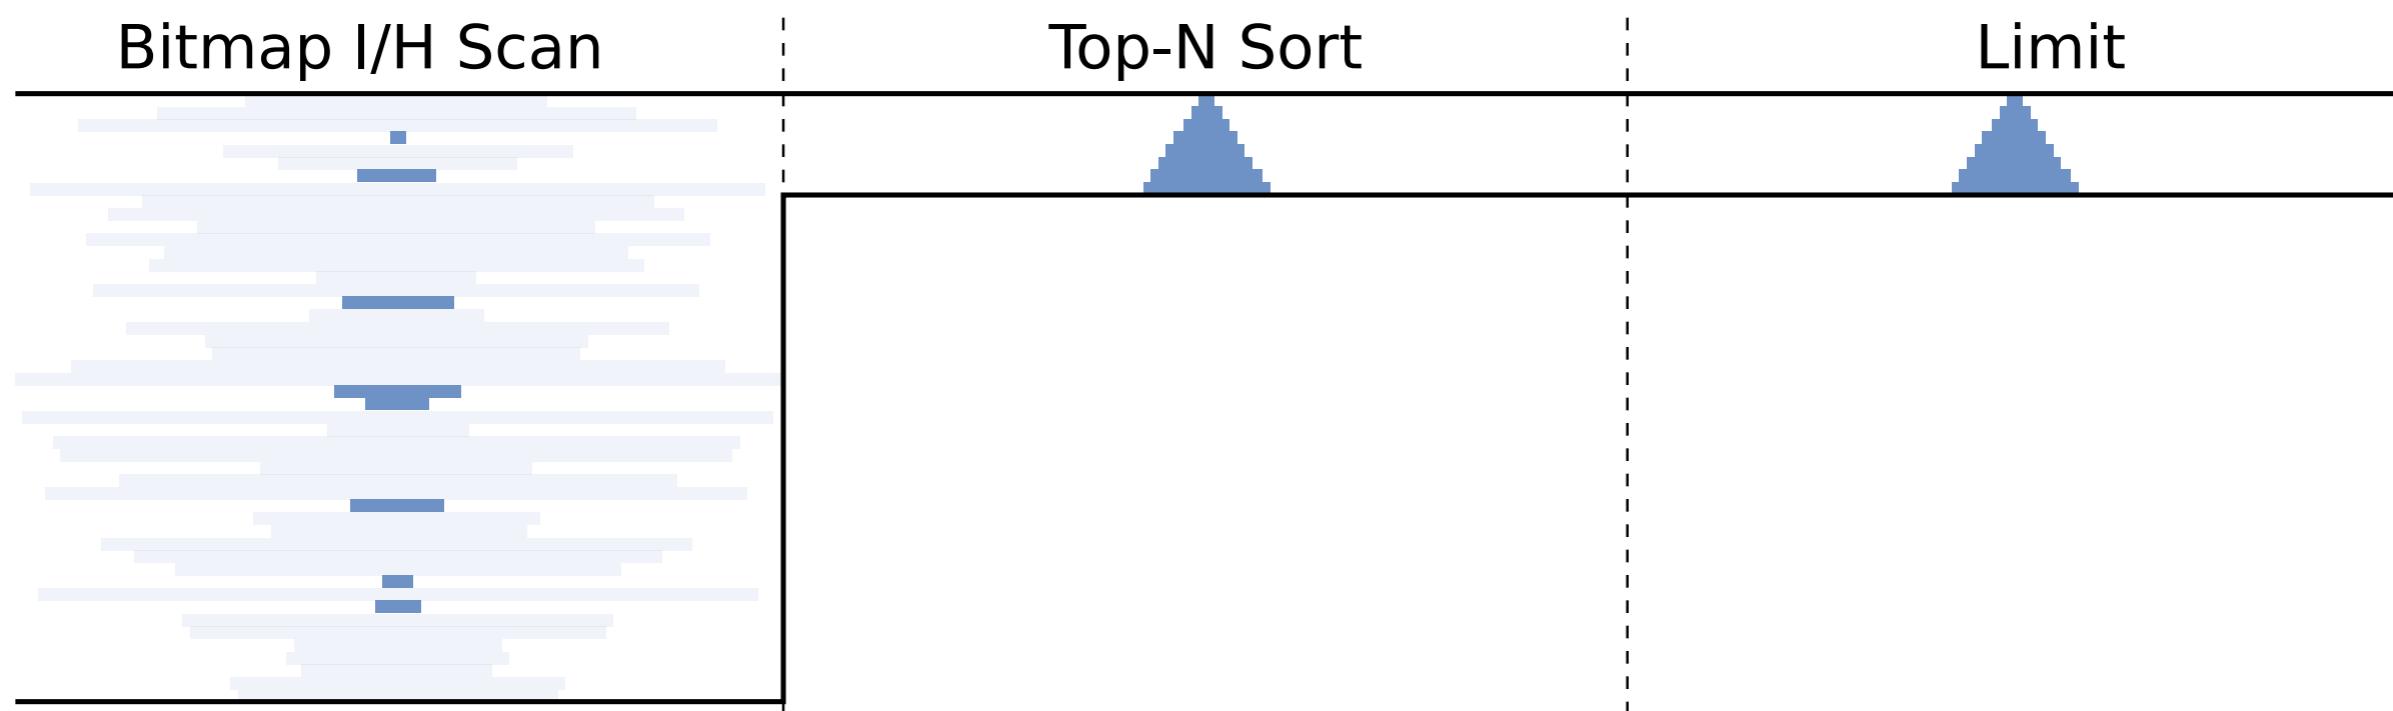
-> Bitmap Heap Scan (**actual rows=9970**)

Rows Removed by Filter: 30 (new in 9.2)

-> Bitmap Index Scan (**actual rows=10000**)

Index Cond: (topic = 1234)

Seek Method w/o Index for order by



Limit (actual rows=10)

-> Sort (**actual rows=10**)

Sort Method: **top-N heapsort** Memory: 18kB

-> Bitmap Heap Scan (**actual rows=10**)

Rows Removed by Filter: 9990 (new in 9.2)

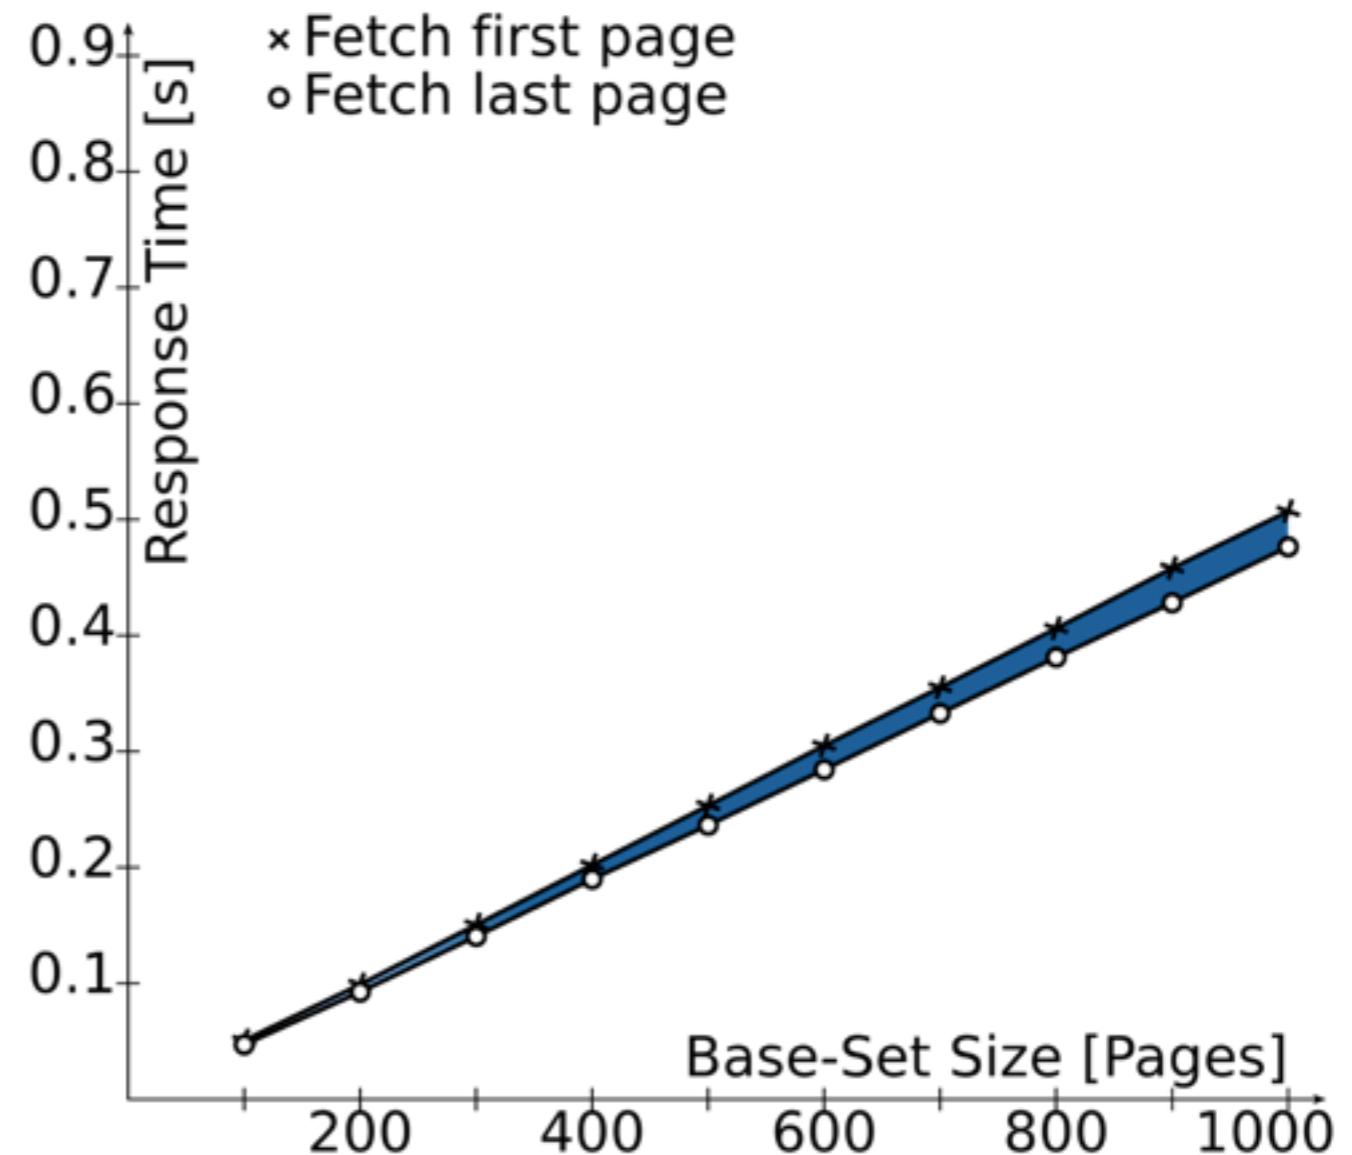
-> Bitmap Index Scan (**actual rows=10000**)

Index Cond: (topic = 1234)

Seek Method w/o Index for order by

Always needs to retrieve the full base set, but the top-n sort buffer needs to hold only 10 rows.

The response time remains the same even when browsing to the last page. And the memory footprint is very low!



Seek-Method with Optimal Index

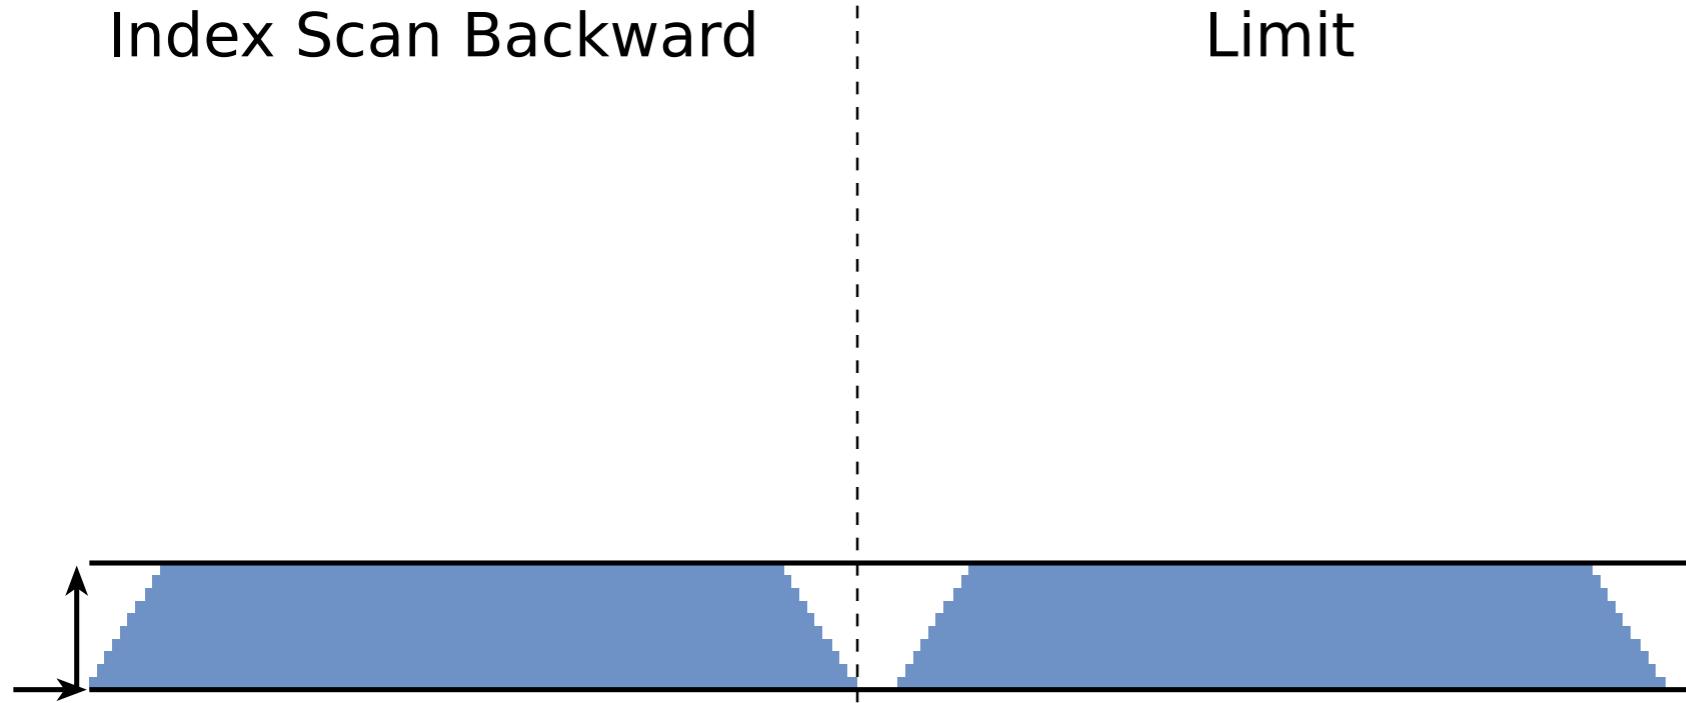
```
select *
  from news
 where topic = 1234
   and (date, id) < (prev_date, prev_id)
 order by date desc, id desc
 limit 10;
```

create index .. on news (topic, date, id);

Seek Method with index for order by

Index Scan Backward

Limit



Limit (**actual rows=10**)

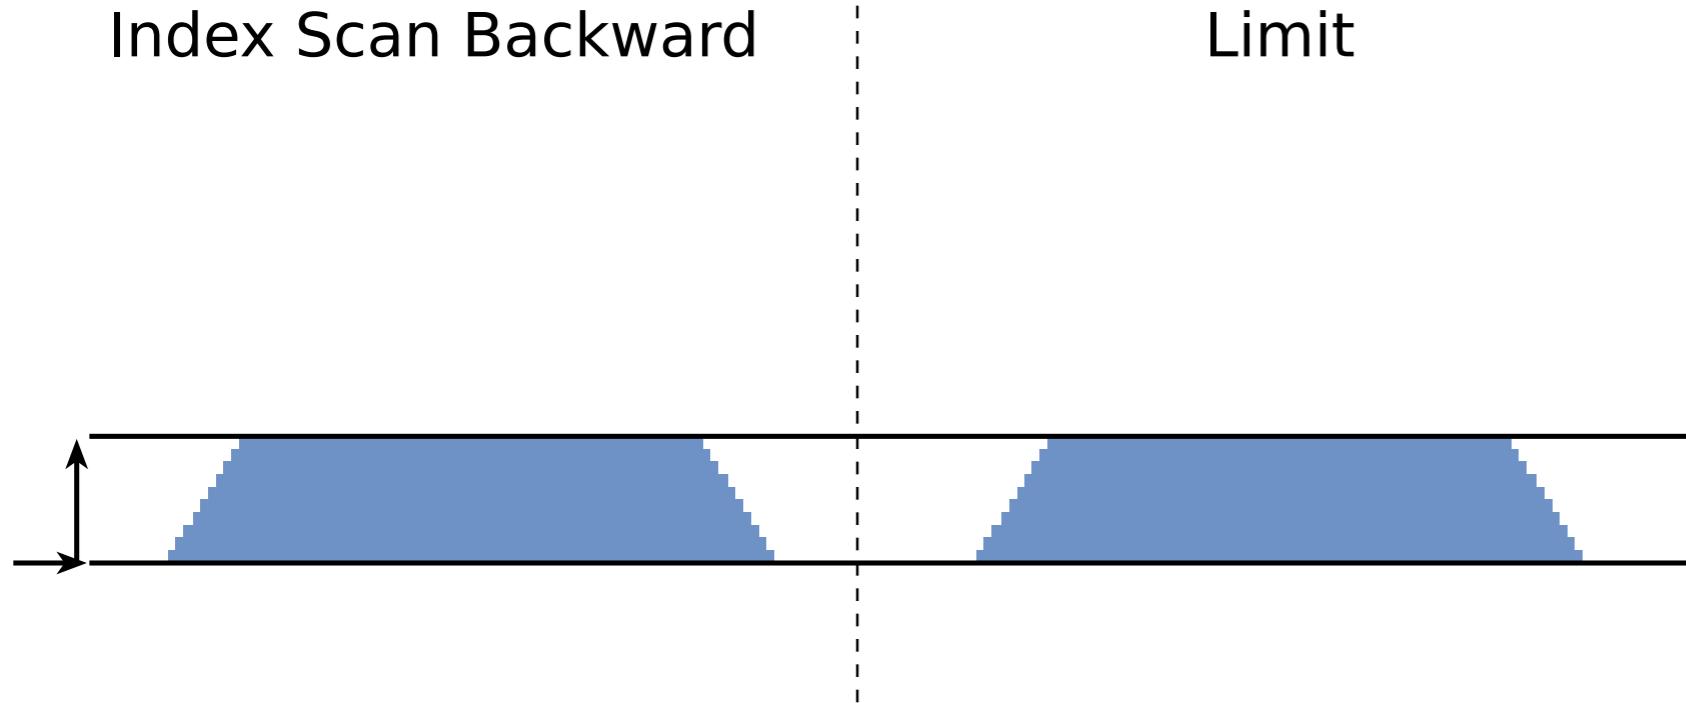
-> Index Scan Backward (**actual rows=10**)

Index Cond: (topic = 1234)

Seek Method with index for order by

Index Scan Backward

Limit



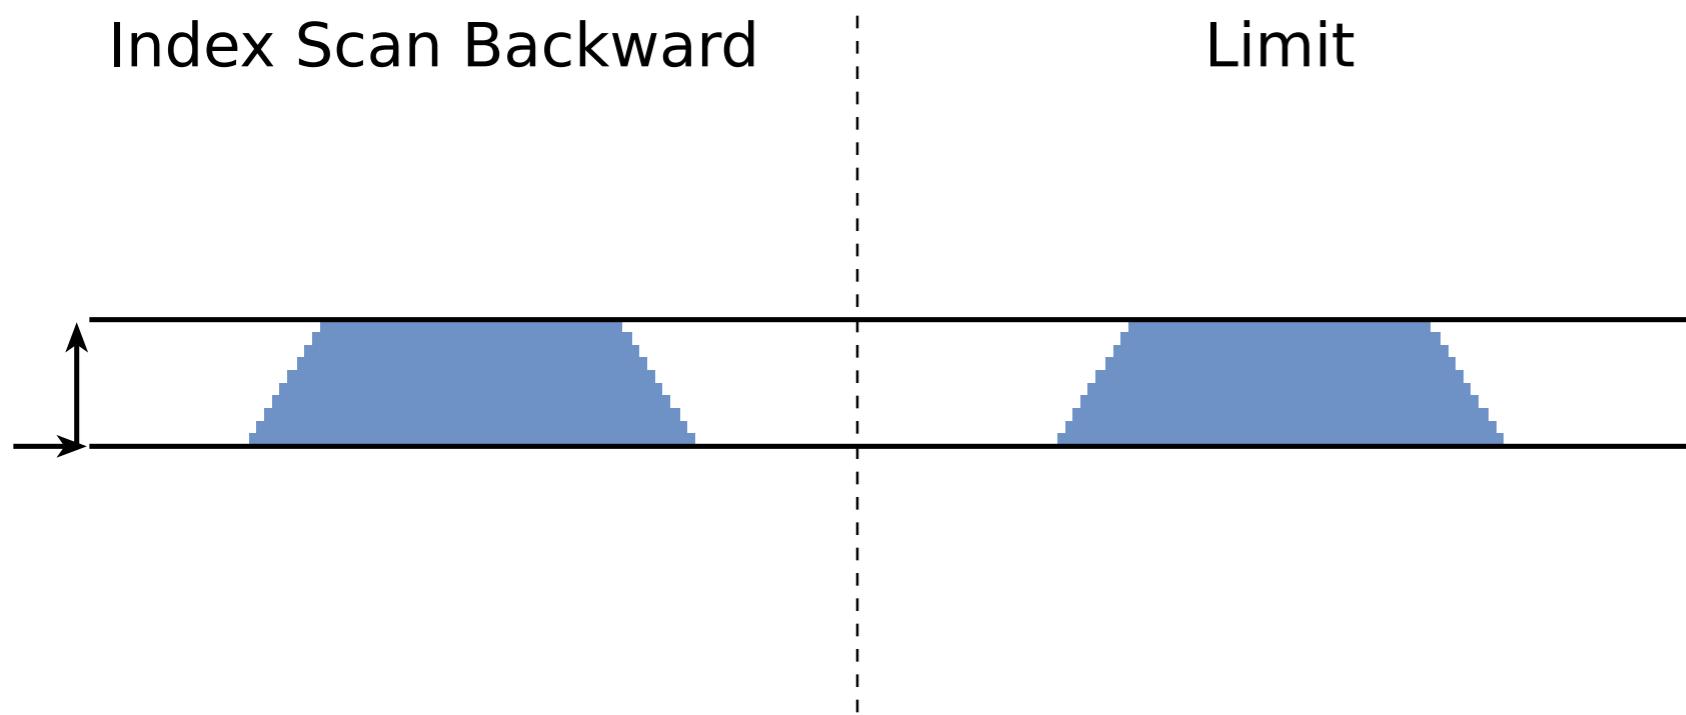
Limit (**actual rows=10**)

-> Index Scan Backward (**actual rows=10**)

Index Cond: ((topic = 1234)

AND (ROW(dt, id) < ROW('...', 23456)))

Seek Method with index for order by



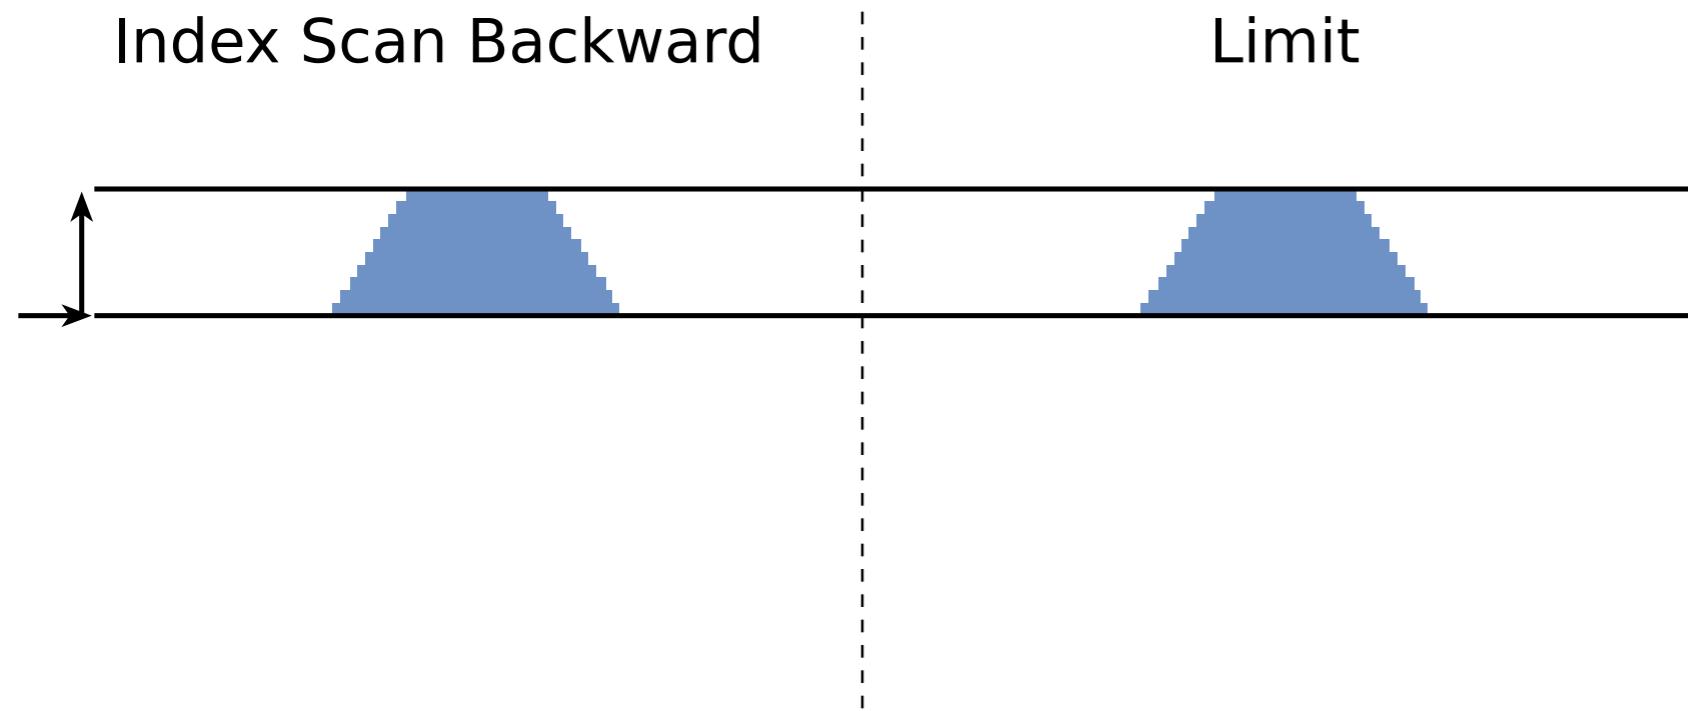
Limit (**actual rows=10**)

-> Index Scan Backward (**actual rows=10**)

Index Cond: ((topic = 1234)

AND (ROW(dt, id) < ROW('...', 34567)))

Seek Method with index for order by



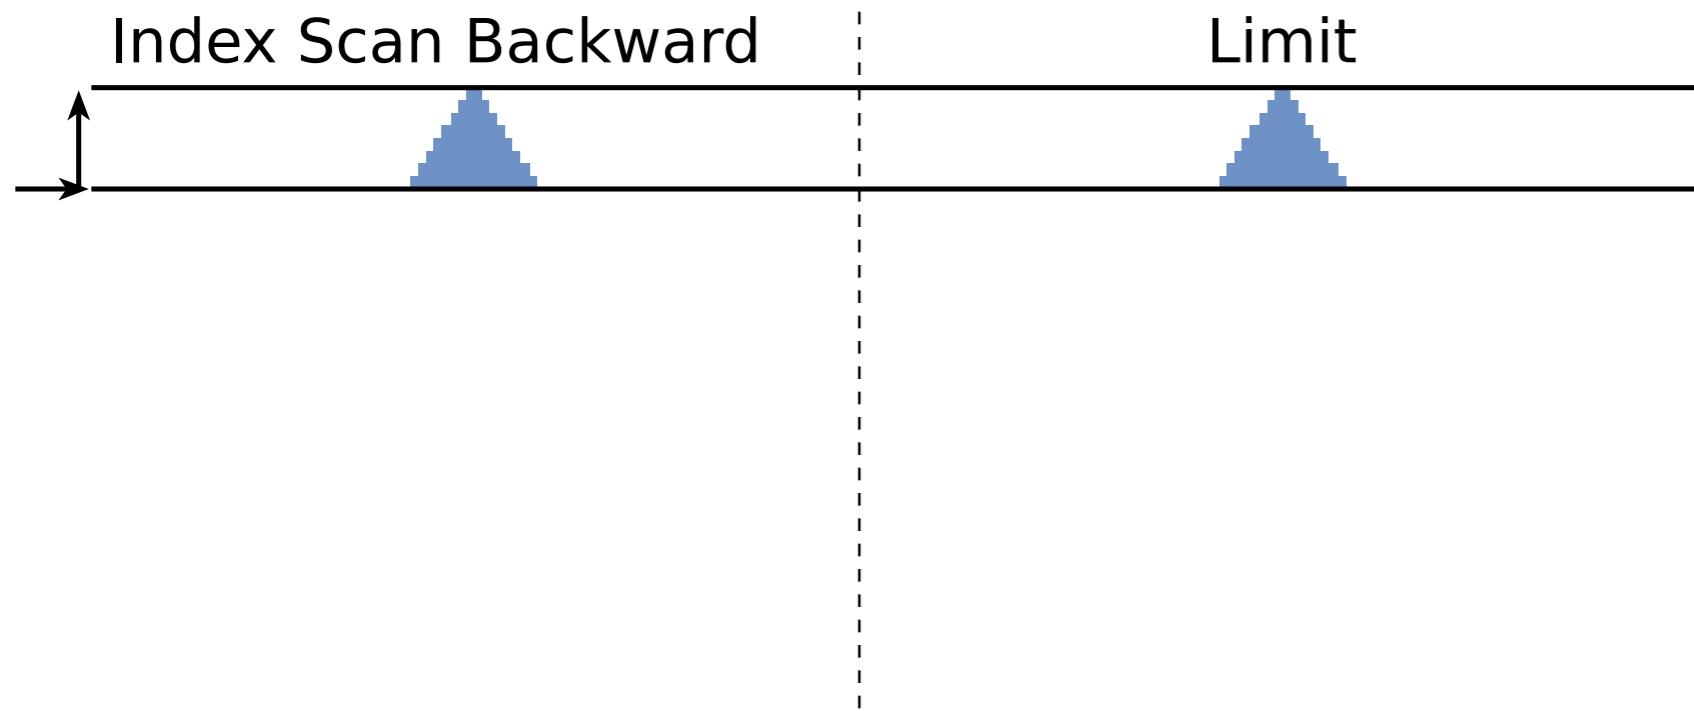
Limit (**actual rows=10**)

-> Index Scan Backward (**actual rows=10**)

Index Cond: ((topic = 1234)

AND (ROW(dt, id) < ROW('...', 45678)))

Seek Method with index for order by



Limit (**actual rows=10**)

-> Index Scan Backward (**actual rows=10**)

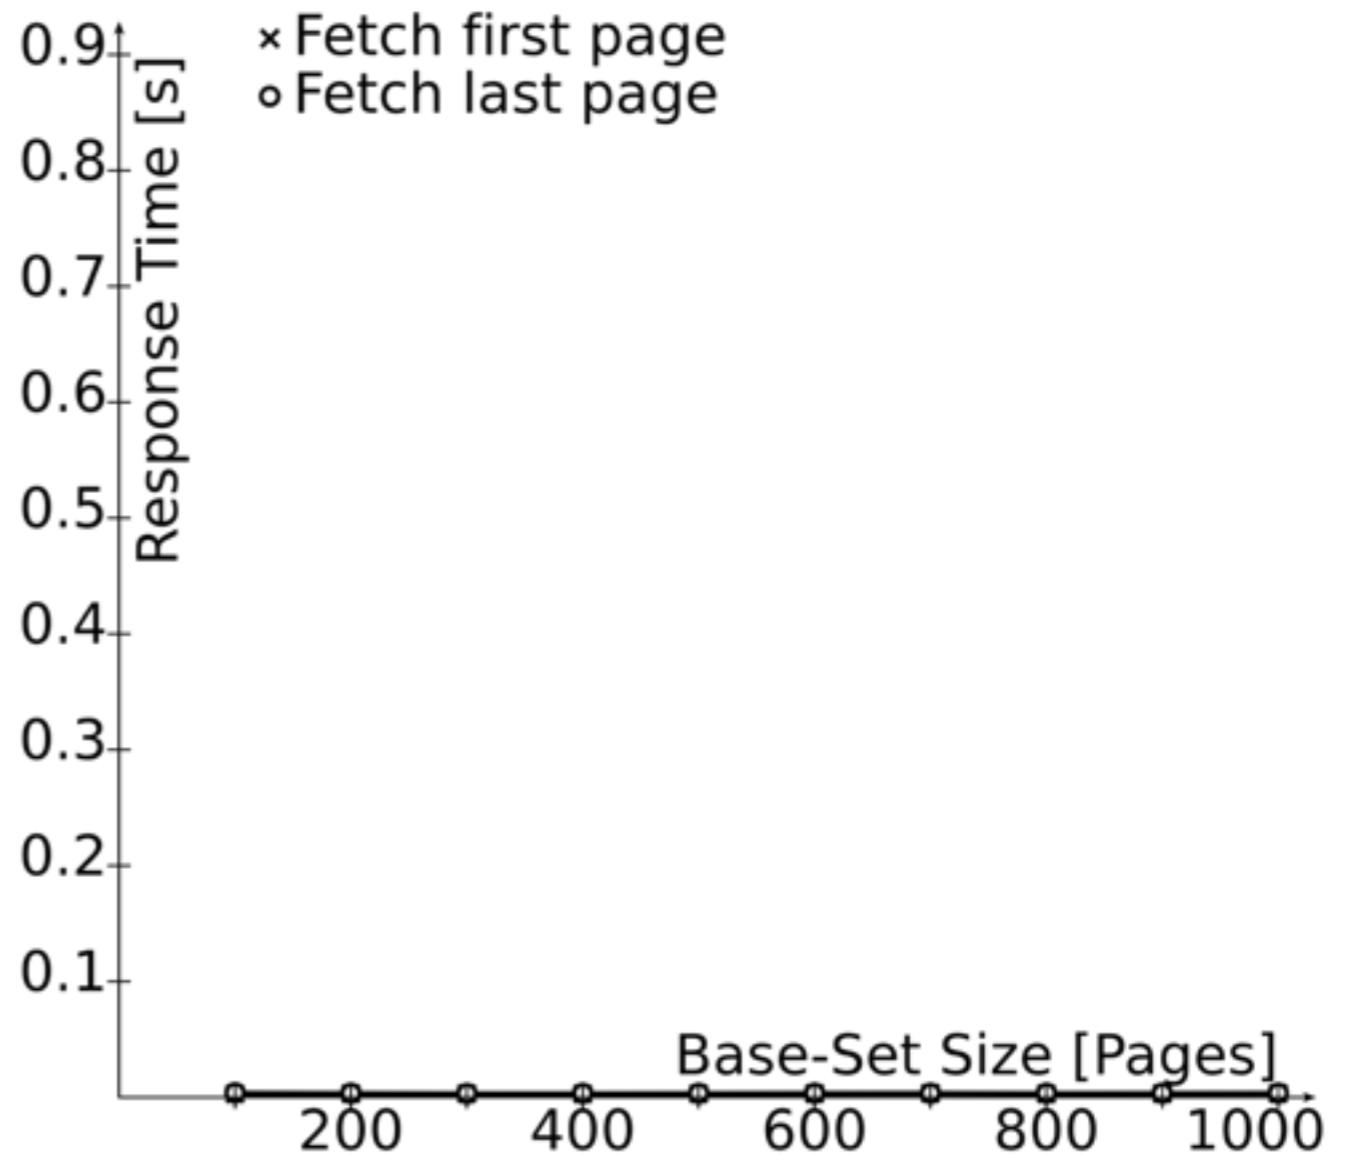
Index Cond: ((topic = 1234)

AND (ROW(dt, id) < ROW('...', 56789)))

Seek Method with index for order by

Successively browsing back doesn't slow down.

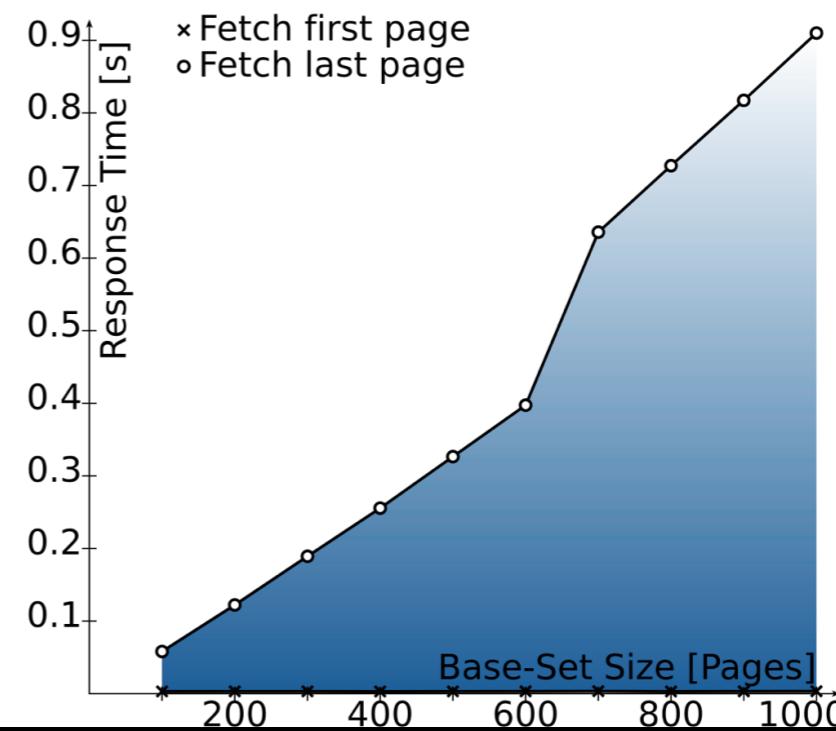
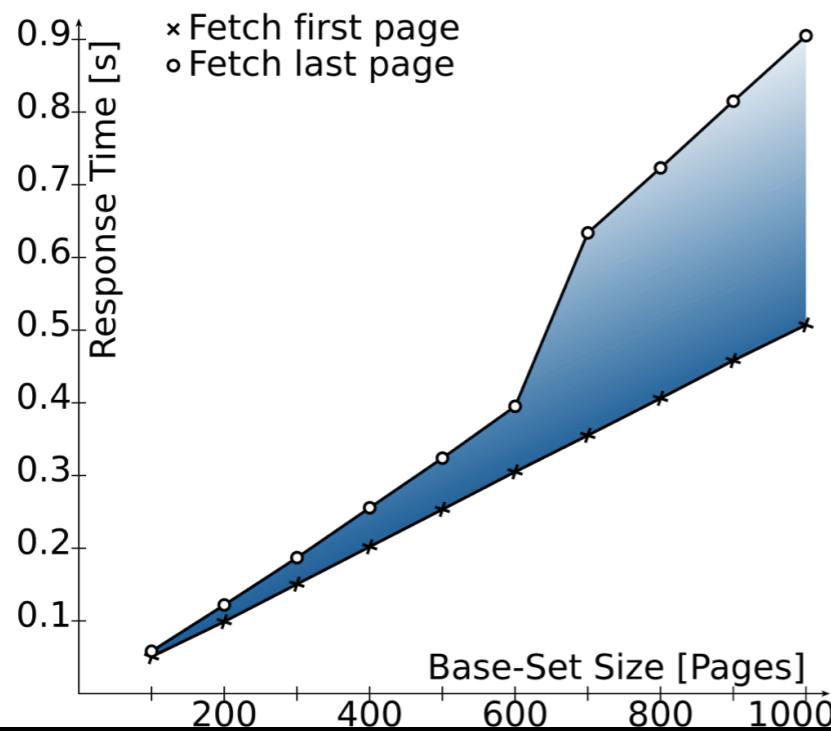
Neither the size of the base set^(*) nor the fetched page number affects the response time.



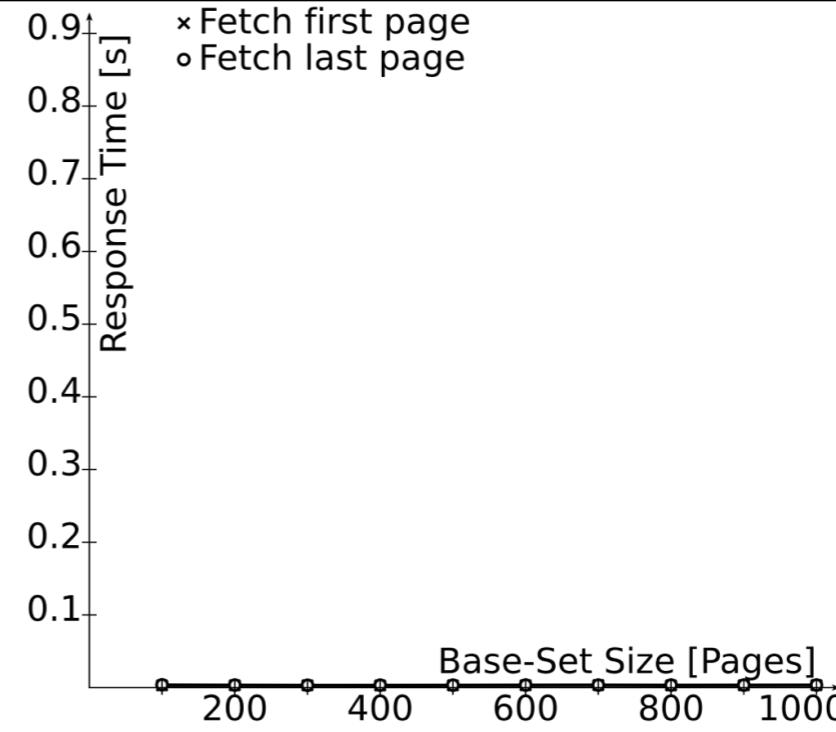
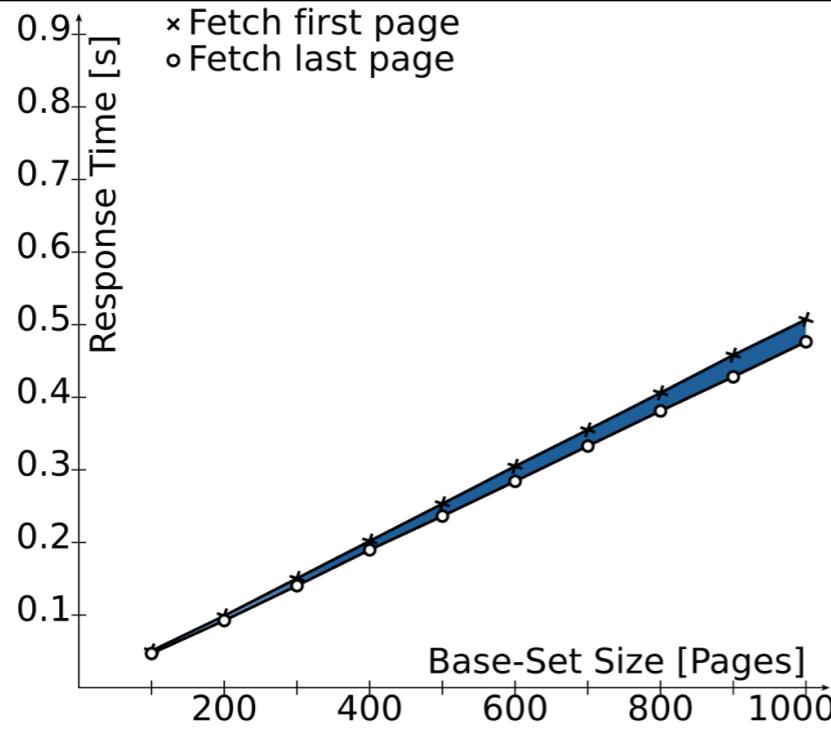
(*) the index tree depth still affects the response time.

Comparison

Offset



Seek



W/O index for `order by`

With index for `order by`

Too good to be true?

The Seek Method has serious limitations

- ▶ You cannot directly navigate to arbitrary pages
 - ▶ because you need the values from the previous page
- ▶ Bi-directional navigation is possible but tedious
 - ▶ you need to reverse the order by direction and RV comparison
- ▶ Works best with full row values support
 - ▶ Workaround is possible, but ugly and less performant
 - ▶ Framework support?
 - ▶ jOOQ 3.3 introduced native support for the Seek-Method
 - ▶ order_query for Ruby by @glebm: http://github.com/glebm/order_query

A Perfect Match for Infinite Scrolling

The “Infinite Scrolling” UI doesn’t need to ...

- ▶ navigate to arbitrary pages
 - ▶ there are no buttons
- ▶ Browse backwards
 - ▶ previous pages are still in the browser
- ▶ Stable results
 - ▶ New rows don’t affect the result
 - No need for de-duplication in JS

SQL Performance Tips @SQLPerfTips 27 Dec
Creating the best index for a WHERE clause is *not* trivial. Just look at the contents of this chapter: is.gd/l0pdzh
[Expand](#)

SQL Performance Tips @SQLPerfTips 27 Dec
@givenilux Sure. Start by filling in this form: sql-performance-explained.com/book/sql-perfo...
[View conversation](#)

SQL Performance Tips @SQLPerfTips 27 Dec
@givenilux Paperback €29.95, PDF €9.95 both together: €34.95.
Shipping would take about 7-10 days + potential delays caused by customs.
[View conversation](#)

SQL Performance Tips @SQLPerfTips 27 Dec
Thanks! Next shipping is tomorrow :) RT @jdumlin Finally bought SQL Performance Explained via @SQLPerfTips #YourTwitterFeedWorked :)
[View conversation](#)

SQL Performance Tips @SQLPerfTips 27 Dec
Have you ever seen a query that was slow although it used an index? Here is why. is.gd/uNLLoo — Disponible en français
[Expand](#)

Also a Perfect Match for PostgreSQL

row values
support matrix

	MySQL	Oracle	PostgreSQL	SQLite	SQL Server
Supported in where clause	✓	✓	✓	✗	✗
Ranges supported (<, >)	✓	✗	✓	✗	✗
Optimal index usage	✗	✓	✓	✗	✗

↑
Popular
Advanced

order by
support matrix

	MySQL	Oracle	PostgreSQL	SQL Server
Read index backwards	✓	✓	✓	✓
Order by ASC/DESC	✓	✓	✓	✓
Index ASC/DESC	✗	✓	✓	✓
Order by NULLS FIRST/LAST	✗	✓	✓	✗
Default NULLS order	First	Last	Last	First
Index NULLS FIRST/LAST	✗	✗	✓	✗

↑
Popular
Advanced

About Markus Winand

Tuning developers for
high SQL performance

Training & co (one-man show):
winand.at

Geeky blog:
use-the-index-luke.com

Author of:
SQL Performance Explained

